


# Computer Architecture

---

The usual state of affairs is a little weird:



- A computer's disk contains some stuff:
  - some of it represents *data*
  - some of it represents *code*
- Before the CPU can use the stuff, it has to stick it in RAM

□

*Q: Why not just run code and use data straight off the disk?*

**A:**


*Q: Okay, then, why not just store everything in RAM all the time?*

**A:**


□

# Speed vs. Persistence

*Speed (in bytes per second)*



*Amount of time data will persist after the power is shut off*



## Specifically...

---

### My PC

---

<i>Pentium II</i>	<i>333MHz</i>
32× CD-ROM	75 ms access
1.44 MB floppy	90 ms access
4.5 GB hard disk	12 ms access
128 MB RAM	60 ns access

---

Given these specs, how many random bytes can you get from each device in *one second*?

	<i>CD</i>	<i>floppy</i>	<i>hard disk</i>	<i>RAM</i>
<i>bytes/sec</i>	<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>

How long would it take to get 10 megs at random off each device?

	<i>CD</i>	<i>floppy</i>	<i>hard disk</i>	<i>RAM</i>
<i>time</i>	<input type="text"/>	<input type="text"/>	<input type="text"/>	<input type="text"/>

These tables explain why hard disks and CD-ROM drives are not used for computer storage: *they're too slow!*

## But Seriously Folks...

---

- The program `memio` (written in C) grabs one byte at a time from random locations within a 100MB space, 10,485,760 times (10 megatimes).
  - *Total time for the move:  $< 1/2$  s*  
(that's about what we'd expect)
- The program `fileio` (written in C) grabs one byte at a time from random locations in a 100MB file.
  - *Total time to move 10,485 random bytes:  $\sim 90$  s*
  - *Total time to move 104,857 random bytes:  $\sim 480$  s*
  - *Total time to move 1,048,576 random bytes:  $\sim 4000$  s*
  - *Total time to move 10,485,760 random bytes:  $\sim ?$  s*  
(dunno... I got bored at this point)
- The program `seqfile` (written in C) grabs 10,485,760 bytes one at a time *in order* starting at a random location within a 100MB file
  - *Total time for the move:  $\sim 60$  s*  
(that's about 1500× faster than random)

# What to Do

---

1. ***maximize*** use of RAM for program operations
  - if you know that you're going to be bouncing around some big block of data, stick it in RAM first
2. ***minimize*** number of disk accesses
  - going to disk once for 1,000 bytes is better than going to disk 1,000 times for 1 byte.
3. ***organize*** information carefully on disk
  - find a way to organize data so you can jump to the right place on disk instead of having to search through large files.
4. ***optimize*** systems to take advantage of 1 - 3
  - write your code and data structures to take advantage of good memory management, disk management and file organization.

*This course is where we find out how to do 1 - 4*

# What We'll Do

---

1. Housekeeping, introduction, motivation
2. Basic file structures, basic operations on files
3. Secondary storage, physical storage devices
4. Managing secondary storage
5. The primary/secondary storage interface, managing primary storage
6. File compression
7. Storage optimization
8. Sorting, searching, merging
9. Indexing
10. B-trees
11. B+-trees
12. Hashing, extendible hashing