

# Record-Based Processing

---

*... or: "How to do stuff to files made up of records".*

## ***Topic***

## ***Folk & Zoellick***

- Old Hat & New Hat
- Basic Record Operations
- Searching
  - Direct Access § 4.2.4
  - Sequential Search §§ 4.2.2, 5.3.3
  - Binary Search §§ 5.3.2, 5.3.3
- Cosequential Processing
  - Matching § 7.1.1
  - Merging § 7.1.2
- Batch Updating
- Space Reclamation § 5.2

# The Key

---

## ***Key***

some subset of fields

## ***Unique Key***

can be used to identify each record in file uniquely

## ***Dataless Key***

contains no information specific to the record



*What's the point of avoiding real data in keys?*

- real data (e.g. Last Name, First Name, etc.) is rarely guaranteed unique
- real data has a tendency to change

# Order and Disorder

---

The kinds of operations that can be performed often depend on whether the records in a file are ordered or unordered.

If the records are ordered, the file is considered *sorted*.

If the records are unordered, the file is considered *unsorted*.



File  $F$  contains records  $r_1 \dots r_n$  having keys  $k_1 \dots k_n$ .

***F is sorted, increasing:***

$$k_1 < k_2 < \dots < k_{n-1} < k_n$$

***F is sorted, non-decreasing:***

$$k_1 \leq k_2 \leq \dots \leq k_{n-1} \leq k_n$$

***F is sorted, decreasing:***

$$k_1 > k_2 > \dots > k_{n-1} > k_n$$

***F is sorted, non-increasing:***

$$k_1 \geq k_2 \geq \dots \geq k_{n-1} \geq k_n$$

# What Sort of File?

---

**Q:** Are these files sorted? How?

5	10	1	a	2	charlie
4	11	12	ab	20	bravo
3	12	12	ab	-2	alpha
2	AA	121	aba	0005	Foxtrot
1	AB	123	abc	01	
0	AC	2	b		
-1		22	bb		
		241	bdc		
		34	cd		

**A:**

# Big O

---

Complexity analysis allows us to compare the efficiency of very different algorithms. If an algorithm is doing some *thing* with  $n$  items, the "big O" notation tells us how the performance of the algorithm varies as  $n$  varies.

```
for(i = 0; i < n; i++)  
    for(j = 0; j < n; j++)  
        ch = getc(infil);
```

Let's say every read (`getc()`) takes 1 ms.

**Q:** If  $n$  is 4, how many milliseconds does the program spend reading?

**A:**

**Q:** What if we increase  $n$  to 5? ... to 6?

**A:**

**Q:** What is the "complexity" (in big O notation) of our program?

**A:**

**Q:** Is the number of reads a fair measure of complexity of this program?

**A:**

# Basic Record Operations

---

Things to do to Records:

- *add* new records to a file
  - where should the record be added? (beginning? middle? end?)
  - if added in the middle or at the beginning, what happens to the keys of the rest of the records?
  - if added in the middle or at the beginning, is there room available
- *modify* existing records in a file
  - does the modification change the length of the record?
  - if the length changes, what happens to the keys of the rest of the records?
  - does the key change?
- *delete* an existing record from a file
  - do we erase the record or simply mark it as deleted?
  - if the record deleted is in the middle of the file, do we have to shuffle all the remaining records around?
  - do the keys of other records change?
  - what do we do with the space available after deleting a record?

If you can figure out how to *add* and *delete* properly, *modify* doesn't really pose any new problems.

## Searching: Direct Access

---

### ***Searching:***

Looking for a particular record  $r_i$  in a file.

### ***Direct Access:***

Seeking directly to the beginning of  $r_i$  without having to read any other records in the file.

### ***If we have fixed length records...***

knowing the record number  $i$  of  $r_i$  allows us to seek directly to the beginning of the record ( $i$  is called the *relative record number* or *RRN*).

### ***If we have variable length records...***

we must know the *byte offset* of  $r_i$  to seek directly to the beginning of the record.

***Q:*** What kind of *key* is *RRN*? What else?

***A:***



*Hmmmm... if only there were a way to know the byte offset of every record in a file....*

## Searching: Sequential Search

---

RRN and byte offset are decent keys: they're unique, they're dataless. But what if we don't know the RRN (or byte offset)?

For example, say we want to find the record for *John Gruden* in this file:

```
Olausson, Fredrik, ana, D, 72, 198
Pushor, Jamie, ana, D, 75, 218
Haller, Kevin, ana, D, 74, 195
...
Matteau, Stephane, was, L, 76, 220
Friesen, Jeff, was, L, 72, 200
Rathje, Mike, was, D, 77, 230
```

We have no choice but to read through all 1,294 records in the file, record-by-record, comparing the first field of each record to *Gruden* and the second field to *John*.

**Q:** How many records will we have to read to find Gruden (best case, worst case, average case)?

**A:**

**Q:** What if Gruden is not in the file?

**A:**

**Q:** What is the complexity (big O) of this search?

**A:**



## Searching: Sequential Search (cont.)

---

But what if that file were *sorted* on the first and second fields?

```
Aalto, Antti, ana, C, 74, 195
Abid, Ramzi, col, L, 74, 195
Abrahamsson, Elias, bos, D, 75, 240
.
.
.
Zubrus, Dainius, phi, R, 75, 220
Zultek, Matt, la, L, 75, 218
Zyuzin, Andrei, sjo, D, 73, 200
```

**Q:** What is the complexity of sequential search through the sorted file?

**A:**

**Q:** What if Gruden is not in *this* file?

**A:**

□

*But we can do better...*

Since the file is sorted, we can look at any record in the file and know immediately if what we're looking for comes before the record or after the record. By looking at the middle record in the file, we can cut the search space in half with just one read!

## Searching: Binary Search

---

```
set low to 0; set high to 1,294

while high ≥ low
    set middle to (high + low) / 2
    read record middle

    if key(middle) = key2find
        return(middle)!
    else if key(middle) < key2find
        set low to middle + 1
    else
        set high to middle - 1
endwhile

return(BUST)
```



**Q:** What is the complexity of binary search through the sorted file?

**A:**

**Q:** What if Gruden is not in the file?

**A:**

**Q:** What if we want to add a record to the file?

**A:**

# Cosequential Processing

Sometimes we have two files (or more) that share some common information. Often we want to find records in the two files that have the same key value.

For example, let's say we have two files: one lists the NHL Eastern Conference leading Ottawa Senators, and the other lists some Scandinavians who have never been in my kitchen.

- *How could we process the records in these two files to find out which Scandinavian Senators have never been in my kitchen?*
- *In fact, none of the Senators (Scandinavian or otherwise) have ever been in my kitchen. How could we process the records in these two files to generate a single complete ordered list?*

Alfredsson	Ahonen
Arvedson	Alfredsson
Berg	Arvedson
Bonk	Björklund
Dackell	Dackell
Gardiner	Doberhof
Gruden	Glad
Hossa	Johansson
Johansson	Jurvanen
Kravchuk	Konestabo
Laukkanen	Laukkanen
Martins	Osnes
McEachern	Puukko
Murray	Salo
Phillips	Trygg
Pitlick	Uggla
Prospal	
Redden	
Rhodes	
Salo	
Sarault	
Traverse	
Tugnutt	
VanAllen	
Yashin	
York	

# Matching

---

*Matching* is the coordinated processing of multiple files to find only records that have the same key in all files (*intersection*).

Here's how it works for two files:

```
open both files  $F_A$   $F_B$ 

get  $k_A$  from  $F_A$ 
get  $k_B$  from  $F_B$ 

while records remain in  $F_A$  and  $F_B$ 
     $k_A < k_B$ : get next  $k_A$  from  $F_A$ 
     $k_A > k_B$ : get next  $k_B$  from  $F_B$ 
     $k_A = k_B$ : output  $k_A$ 
                  get next  $k_A$  from  $F_A$ 
                  get next  $k_B$  from  $F_B$ 
endwhile

close files
```

**Q:** What happens if there is an error in the ordering of a file?

**A:**

# Merging

---

*Merging* is the coordinated processing of multiple files to find the *union* of records in all files (no duplicates).

Here's how it works for two files:

```
open both files  $F_A$   $F_B$ 

get  $k_A$  from  $F_A$ 
get  $k_B$  from  $F_B$ 

while records remain in  $F_A$  or  $F_B$ 
     $k_A < k_B$ : output  $k_A$ 
                get next  $k_A$  from  $F_A$ 
     $k_A > k_B$ : output  $k_B$ 
                get next  $k_B$  from  $F_B$ 
     $k_A = k_B$ : output  $k_A$ 
                get next  $k_A$  from  $F_A$ 
                get next  $k_B$  from  $F_B$ 
endwhile

close files

/* Note: when getting the next  $k_x$  from file  $F_x$ ,
   if  $F_x$  is at end of file, set  $k_x$  to some
   "out-of-range" high value. */
```

**Q:** Wow, isn't this almost the same algorithm as the Matching algorithm?

# Batch Updating

---

When there are frequent changes (additions, modifications, deletions) to records in a data file, and instantaneous updating is not required, it is often more practical to "save up" the changes in a separate file (called a *transaction* file) until there are enough to warrant updating the data file.

## *Data File*

```
Alfredsson, Daniel, R, 71, 194
Bonk, Radek, C, 75, 224
Hossa, Marian, L, 73, 194
Murray, Chris, R, 74, 215
Neckar, Stan, D, 73, 212
Phillips, Chris, D, 74, 218
Yashin, Alexei, C, 75, 228
```

## *Transaction File*

```
98/11/27: del, Neckar, Stan
98/11/27: add, Berg, Bill, L, 73, 205
99/01/01: mod, Murray, Chris, ?, ?, 222
99/02/27: add, Barker, Ken, D, 70, ?
```

We want to use the *transaction file* to update the *data file*. But if we leave the transaction file in its current order, we'll have to do a search through the data file for every record in the transaction file.

*There must be a better way...*

## Batch Updating (cont.)

---

If we sort the transaction file on the same key as the data file, we can use *cosequential processing* to update the data file in one loop!

### *Transaction File*

```
99/02/27: add, Barker, Ken, D, 70, ?  
98/11/27: add, Berg, Bill, L, 73, 205  
99/01/01: mod, Murray, Chris, ?, ?, 222  
98/11/27: del, Neckar, Stan
```

And let's repeat the data file just for fun.

### *Data File*

```
Alfredsson, Daniel, R, 71, 194  
Bonk, Radek, C, 75, 224  
Hossa, Marian, L, 73, 194  
Murray, Chris, R, 74, 215  
Neckar, Stan, D, 73, 212  
Phillips, Chris, D, 74, 218  
Yashin, Alexei, C, 75, 228
```

# Batch Updating Algorithm

---

The algorithm for updating might look like this:

```
open both files  $F_D$   $F_T$ 

get  $k_D$  from  $F_D$ 
get  $k_T$  from  $F_T$ 

while records remain in  $F_D$  or  $F_T$ 
     $k_D < k_T$ : output  $r_D$ , get next  $k_D$  from  $F_D$ 
     $k_D > k_T$ : if action = add
        output  $r_T$ , get next  $k_T$  from  $F_T$ 
    else error
    endif
     $k_D = k_T$ : if action = mod
        output  $r_D$  with new field values
    else if action = del
        do nothing
    else error
    endif

    get next  $k_D$  from  $F_D$ 
    get next  $k_T$  from  $F_T$ 
endwhile

close files

/* Again: when getting the next  $k_x$  from file  $F_x$ ,
   if  $F_x$  is at end of file, set  $k_x$  to some
   "out-of-range" high value (like 'ZZZZZZ'). */
```



## Not Batch Updating

---

The cosequential processing algorithms do exhaustive sequential reading of multiple files. For example, in the batch updating scenario, if there are  $N$  records in the data file and  $M$  records in the transaction file, the algorithm is  $O(M + N)$  (or  $O(\max(M, N))$ ).

What if *instantaneous* updating is required? That is, every time a transaction takes place, the data file must be updated immediately.

**Q:** What is the complexity of instantaneous updating of a data file with  $N$  records with  $M$  transactions?

**A:**



The updating we've seen writes a new file with the updated information. That's kind of a waste if we're only changing one record! Can't we just make the changes directly to the data file?

## Modifying a File Directly

---

Suppose we have to make some changes to the following file and we don't care about maintaining order in the file.

```
1  Olausson, Fredrik, ana, D, 72, 198
2  Pushor, Jamie, ana, D, 75, 218
3  Haller, Kevin, ana, D, 74, 195
4  Trnka, Pavel, ana, D, 75, 200
...
n-2 Matteau, Stephane, was, L, 76, 220
n-1 Friesen, Jeff, was, L, 72, 200
n   Rathje, Mike, was, D, 77, 230
```

Let's say David Oliver gets called up from the minors. We need to add the record

Oliver, David, ott, L, 72, 190

```
1  Olausson, Fredrik, ana, D, 72, 198
2  Pushor, Jamie, ana, D, 75, 218
3  Haller, Kevin, ana, D, 74, 195
4  Trnka, Pavel, ana, D, 75, 200
...
n-2 Matteau, Stephane, was, L, 76, 220
n-1 Friesen, Jeff, was, L, 72, 200
n   Rathje, Mike, was, D, 77, 230
n+1 Oliver, David, ott, L, 72, 190
```

Adding a record to our file was pretty simple. What about deleting a record?

## Storage Compaction

---

Let's say Kevin Haller retires at the ripe old age of 28. Instead of writing a brand new file with Haller's record deleted, we can just mark the record with a special mark that means: "this record is deleted". Any software using the file will just skip over records marked as deleted.

```
1   Olausson, Fredrik, ana, D, 72, 198
2   Pushor, Jamie, ana, D, 75, 218
3   ~ Haller, Kevin, ana, D, 74, 195
4   Trnka, Pavel, ana, D, 75, 200
...
n-2  Matteau, Stephane, was, L, 76, 220
n-1  Friesen, Jeff, was, L, 72, 200
n     Rathje, Mike, was, D, 77, 230
n+1  Oliver, David, ott, L, 72, 190
```

Every once in a while, we can run a special *compactor* program that writes out a new file omitting all the deleted records (this is known as *space reclamation* by *storage compaction*).

Deleting records in this manner is simple and fast. But there's another advantage as well...

**Q:** What's the other advantage?

**A:**

**Q:** What's the *other* other advantage?

**A:**

## Dynamic Reclamation

---

The problem with reusing the space left by deleted records is that we have to search sequentially through the file for the available spots.

A better solution would be to use a data structure (let's say, a *stack*) to keep track of which records are deleted.

We store the RRN of the deleted record at the top of the stack in a special record in the file. The file below has no deleted records.

0	
1	Olausson, Fredrik, ana, D, 72, 198
2	Pushor, Jamie, ana, D, 75, 218
3	Haller, Kevin, ana, D, 74, 195
4	Trnka, Pavel, ana, D, 75, 200
5	Barker, Ken, ana, D, 70, ???
...	
n-2	Matteau, Stephane, was, L, 76, 220
n-1	Friesen, Jeff, was, L, 72, 200
n	Rathje, Mike, was, D, 77, 230
n+1	Oliver, David, ott, L, 72, 190

## Dynamic Reclamation (cont.)

---

If we then delete the record with RRN 5, we mark the record as deleted, change the special "next in stack" field to 0, and change the special "top of stack" record to 5.

	<b>5</b>	
1		Olausson, Fredrik, ana, D, 72, 198
2		Pushor, Jamie, ana, D, 75, 218
3		Haller, Kevin, ana, D, 74, 195
4		Trnka, Pavel, ana, D, 75, 200
5	<b>~ 0</b>	<b>Barker, Ken, ana, D, 70, ???</b>
	<b>...</b>	
n-2		Matteau, Stephane, was, L, 76, 220
n-1		Friesen, Jeff, was, L, 72, 200
n		Rathje, Mike, was, D, 77, 230
n+1		Oliver, David, ott, L, 72, 190

## Dynamic Reclamation (cont.)

---

Then if we delete record 3, we mark the record as deleted, change the "next in stack" field to the value of the current "top of stack", and change the "top of stack" to 3.

	<b>3</b>	
1		Olausson, Fredrik, ana, D, 72, 198
2		Pushor, Jamie, ana, D, 75, 218
3	~ 5	<b>Haller, Kevin, ana, D, 74, 195</b>
4		Trnka, Pavel, ana, D, 75, 200
5	~ 0	Barker, Ken, ana, D, 70, ???
...		
n-2		Matteau, Stephane, was, L, 76, 220
n-1		Friesen, Jeff, was, L, 72, 200
n		Rathje, Mike, was, D, 77, 230
n+1		Oliver, David, ott, L, 72, 190

## Dynamic Reclamation: The Reclaiming Part

---

The next time we add a record, we can reclaim deleted space simply by jumping directly to the deleted record at the top of the stack:

```
      5
1      Olausson, Fredrik, ana, D, 72, 198
2      Pushor, Jamie, ana, D, 75, 218
3      Leclerc, Mike, ana, L, 73, 205
4      Trnka, Pavel, ana, D, 75, 200
5 ~ 0  Barker, Ken, ana, D, 70, ???
...
n-2    Matteau, Stephane, was, L, 76, 220
n-1    Friesen, Jeff, was, L, 72, 200
n      Rathje, Mike, was, D, 77, 230
n+1    Oliver, David, ott, L, 72, 190
```

## Food for Thought

---

*What if the records are variable length?*



***Q:*** How do we refer to available records in the stack?

***A:***

***Q:*** What if the space left by the deleted record is not big enough for the record we want to add?

***A:***