

Chapter 12

Concurrency can occur at four levels:

- 1. Machine instruction level**
- 2. High-level language statement level**
- 3. Unit level**
- 4. Program level**

Because there are no language issues in instruction- and program-level concurrency, they are not addressed here

The Evolution of Multiprocessor Architectures

- 1. Late 1950s - One general-purpose processor and one or more special-purpose processors for input and output operations**
- 2. Early 1960s - Multiple complete processors, used for program-level concurrency**
- 3. Mid-1960s - Multiple partial processors, used for instruction-level concurrency**
- 4. Single-Instruction Multiple-Data (SIMD) machines**
The same instruction goes to all processors, each with different data - e.g., *vector processors*

Chapter 12

5. Multiple-Instruction Multiple-Data (MIMD) machines

- Independent processors that can be synchronized (unit-level concurrency)
-

Def: A *thread of control* in a program is the sequence of program points reached as control flows through the program

Categories of Concurrency:

1. *Physical concurrency* - Multiple independent processors
(multiple threads of control)

2. *Logical concurrency* - The appearance of physical concurrency is presented by time-sharing one processor
(software can be designed as if there were multiple threads of control)

- *Coroutines provide only quasiconcurrency*

Chapter 12

Reasons to Study Concurrency

1. It involves a new way of designing software that can be very useful--many real-world situations involve concurrency
2. Computers capable of physical concurrency are now widely used

Fundamentals (for stmt-level concurrency)

Def: A *task* is a program unit that can be in concurrent execution with other program units

- Tasks differ from ordinary subprograms in that:
 1. A task may be implicitly started
 2. When a program unit starts the execution of a task, it is not necessarily suspended
 3. When a task's execution is completed, control may not return to the caller

Def: A task is *disjoint* if it does not communicate with or affect the execution of any other task in the program in any way

Chapter 12

Task communication is necessary for synchronization

- *Task communication can be through:*

- 1. Shared nonlocal variables**
- 2. Parameters**
- 3. Message passing**

- *Kinds of synchronization:*

1. *Cooperation*

- Task A must wait for task B to complete some specific activity before task A can continue its execution
e.g., the producer-consumer problem**

2. *Competition*

- When two or more tasks must use some resource that cannot be simultaneously used
e.g., a shared counter**
- A problem because operations are not atomic**

Chapter 12

- Competition is usually provided by *mutually exclusive access* (methods are discussed later)
- Providing synchronization requires a mechanism for delaying task execution
- Task execution control is maintained by a program called the scheduler, which maps task execution onto available processors
- Tasks can be in one of several different execution states:
 1. New - created but not yet started
 2. Runnable or ready - ready to run but not currently running (no available processor)
 3. Running
 4. Blocked - has been running, but cannot now continue (usually waiting for some event to occur)
 5. Dead - no longer active in any sense

Chapter 12

- **Liveness** is a characteristic that a program unit may or may not have
 - In sequential code, it means the unit will eventually complete its execution
 - In a concurrent environment, a task can easily lose its liveness
- If all tasks in a concurrent environment lose their liveness, it is called *deadlock*
- **Design Issues for Concurrency:**
 1. How is cooperation synchronization provided?
 2. How is competition synchronization provided?
 3. How and when do tasks begin and end execution?
 4. Are tasks statically or dynamically created?

Chapter 12

Methods of Providing Synchronization:

1. Semaphores
2. Monitors
3. Message Passing

1. Semaphores (Dijkstra - 1965)

- **A *semaphore* is a data structure consisting of a counter and a queue for storing task descriptors**
- **Semaphores can be used to implement guards on the code that accesses shared data structures**
- **Semaphores have only two operations, wait and release (originally called P and V by Dijkstra)**
- **Semaphores can be used to provide both competition and cooperation synchronization**

Chapter 12

- Cooperation Synchronization with Semaphores:

- Example: A shared buffer

- The buffer is implemented as an ADT with the operations `DEPOSIT` and `FETCH` as the only ways to access the buffer**

- Use two semaphores for cooperation:**

`emptyspots` and `fullspots`

- The semaphore counters are used to store the numbers of empty spots and full spots in the buffer**

- `DEPOSIT` must first check `emptyspots` to see if there is room in the buffer**

- If there is room, the counter of `emptyspots` is decremented and the value is inserted**

- If there is no room, the caller is stored in the queue of `emptyspots`**

- When `DEPOSIT` is finished, it must increment the counter of `fullspots`**

Chapter 12

- **FETCH must first check `fullspots` to see if there is a value**
- **If there is a full spot, the counter of `fullspots` is decremented and the value is removed**
- **If there are no values in the buffer, the caller must be placed in the queue of `fullspots`**
- **When `FETCH` is finished, it increments the counter of `emptyspots`**
- **The operations of `FETCH` and `DEPOSIT` on the semaphores are accomplished through two semaphore operations named `wait` and `release`**

```
wait(aSemaphore)  
  if aSemaphore s counter > 0 then  
    Decrement aSemaphore s counter  
  else  
    Put the caller in aSemaphore s queue  
    Attempt to transfer control to some  
    ready task  
    (If the task ready queue is empty,  
    deadlock occurs)  
  end
```

Chapter 12

```
release(aSemaphore)  
  if aSemaphore s queue is empty then  
    Increment aSemaphore s counter  
  else  
    Put the calling task in the task ready  
    queue  
    Transfer control to a task from  
    aSemaphore s queue  
  end
```

---> **SHOW Program (p. 500)**

- Competition Synchronization with Semaphores

- A third semaphore, named `access`, is used to control access (competition synchronization)

- The counter of `access` will only have the values 0 and 1

- Such a semaphore is called a *binary semaphore*

---> **SHOW the complete shared buffer example program (p. 501-502)**

- Note that wait and release must be atomic!

Chapter 12

Evaluation of Semaphores:

1. Misuse of semaphores can cause failures in cooperation synchronization
e.g., the buffer will overflow if the wait of `fullspots` is left out
2. Misuse of semaphores can cause failures in competition synchronization
e.g., The program will deadlock if the release of `access` is left out

2. Monitors (Concurrent Pascal, Modula, Mesa)

The idea: encapsulate the shared data and its operations to restrict access

A *monitor* is an abstract data type for shared data

---> SHOW the diagram of monitor buffer operation, Figure 11.2 (p. 505)

- *Example language: Concurrent Pascal*
- Concurrent Pascal is Pascal + classes, processes (tasks), monitors, and the `queue` data type (for semaphores)

Chapter 12

- *Example language: Concurrent Pascal (continued)*
- Processes are types
 - Instances are statically created by declarations
 - An instance is started by `init`, which allocates its local data and begins its execution
- Monitors are also types
 - Form:

`type some_name = monitor (formal parameters)`
shared variables
local procedures
exported procedures (have `entry` in definition)
initialization code
- *Competition Synchronization with Monitors:*
 - Access to the shared data in the monitor is limited by the implementation to a single process at a time; therefore, mutually exclusive access is inherent in the semantic definition of the monitor
 - Multiple calls are queued

Chapter 12

- Cooperation Synchronization with Monitors:

- Cooperation is still required - done with semaphores, using the `queue` data type and the built-in operations, `delay` (similar to `wait`) and `continue` (similar to `release`)
- `delay` takes a queue type parameter; it puts the process that calls it in the specified queue and removes its exclusive access rights to the monitor's data structure
 - Differs from `send` because `delay` always blocks the caller
- `continue` takes a queue type parameter; it disconnects the caller from the monitor, thus freeing the monitor for use by another process. It also takes a process from the parameter queue (if the queue isn't empty) and starts it
 - Differs from `release` because it always has some effect (`release` does nothing if the queue is empty)

---> SHOW `databuf` monitor (p. 506), producer and consumer processes and the program that uses the buffer (p. 506-507)

Chapter 12

Evaluation of monitors:

- Support for competition synchronization is great!
- Support for cooperation synchronization is very similar as with semaphores, so it has the same problems

3. Message Passing

- Message passing is a general model for concurrency
 - It can model both semaphores and monitors
 - It is not just for competition synchronization
- *Central idea:* task communication is like seeing a doctor--most of the time he waits for you or you wait for him, but when you are both ready, you get together, or rendezvous
- In terms of tasks, we need:
 - a. A mechanism to allow a task to indicate when it is willing to accept messages
 - b. Tasks need a way to remember who is waiting to have its message accepted and some fair way of choosing the next message

Chapter 12

Def: When a sender task's message is accepted by a receiver task, the actual message transmission is called a *rendezvous*

- The Ada 83 Message-Passing Model

- Ada tasks have specification and body parts, like packages; the spec has the interface, which is the collection of entry points.**

e.g. task EX is
 entry ENTRY_1 (STUFF : in FLOAT);
end EX;

- The *body* task describes the action that takes place when a rendezvous occurs**
- A task that sends a message is suspended while waiting for the message to be accepted *and* during the rendezvous**
- Entry points in the spec are described with *accept* clauses (message sockets) in the *body***

Chapter 12

- Example of a task body:

```
task body EX is
  begin
    loop
      accept ENTRY_1 (ITEM: in FLOAT) do
        ...
      end;
    end loop;
  end EX;
```

- Semantics:

- a. The task executes to the top of the `accept` clause and waits for a message
- b. During execution of the `accept` clause, the sender is suspended
- c. `accept` parameters can transmit information in either or both directions
- d. Every `accept` clause has an associated queue to store waiting messages

---> **SHOW** rendezvous time lines for the example task (Figure 12.3, p. 511)

Chapter 12

- A task that has `accept` clauses, but no other code is called a **server task** (the example above is a server task)
- A task without `accept` clauses is called an **actor task**
- Example actor task:

```
task WATER_MONITOR; -- specification
task body WATER_MONITOR is -- body
begin
  loop
    if WATER_LEVEL > MAX_LEVEL
      then SOUND_ALARM;
    end if;
    delay 1.0; -- No further execution
               -- for at least 1 second
  end loop;
end WATER_MONITOR;
```

- An actor task can send messages to other tasks
- Note: A sender must know the `entry` name of the receiver, but not vice versa
- Tasks can be either statically or dynamically allocated

Chapter 12

- **Example:**

```
task type TASK_TYPE_1 is ... end;  
type TASK_PTR is access TASK_TYPE_1;  
TASK1 : TASK_TYPE_1;           -- stack dynamic  
TASK_PTR := new TASK_TYPE_1;   -- heap dynamic
```

- **Tasks can have more than one entry point**

- **The specification task has an entry clause for each**

- **The task body has an accept clause for each entry clause, placed in a select clause, which is in a loop**

- **Example task with multiple entries:**

```
task body TASK_EXAMPLE is  
  loop  
    select  
      accept ENTRY_1 (formal params) do  
        ...  
      end ENTRY_1;  
      ...  
    or  
      accept ENTRY_2 (formal params) do  
        ...  
      end ENTRY_2;  
      ...  
    end select;  
  end loop;  
end TASK_EXAMPLE;
```

Chapter 12

- **Semantics of tasks with `select` clauses:**

- If exactly one `entry` queue is nonempty, choose a message from it
- If more than one `entry` queue is nonempty, choose one, nondeterministically, from which to accept a message
- If all are empty, wait
- Extended `accept` clause - code following the clause, but before the next clause
 - Executed concurrently with the caller

- **Cooperation Synchronization with Message Passing**

- Provided by **Guarded** `accept` clauses

- **Example:**

```
when not FULL(BUFFER) =>
  accept DEPOSIT (NEW_VALUE) do
    ...
  end DEPOSIT;
```

Chapter 12

Def: A clause whose guard is true is called *open*.

Def: A clause whose guard is false is called *closed*.

Def: A clause without a guard is always open.

- Semantics of select with guarded accept clauses:

`select` first checks the guards on all clauses

If exactly one is open, its queue is checked for messages

If more than one are open, nondeterministically choose a queue among them to check for messages

If all are closed, it is a runtime error

- A `select` clause can include an `else` clause to avoid the error

- When the `else` clause completes, the loop repeats

Chapter 12

Example of a task with guarded accept clauses:

```
task GAS_STATION_ATTENDANT is
  entry SERVICE_ISLAND (CAR : CAR_TYPE);
  entry GARAGE (CAR : CAR_TYPE);
end GAS_STATION_ATTENDANT;

task body GAS_STATION_ATTENDANT is
begin
  loop
    select
      when GAS_AVAILABLE =>
        accept SERVICE_ISLAND (
          CAR : CAR_TYPE) do
          FILL_WITH_GAS (CAR);
        end SERVICE_ISLAND;
      or
      when GARAGE_AVAILABLE =>
        accept GARAGE (
          CAR : CAR_TYPE) do
          FIX (CAR);
        end GARAGE;
      else
        SLEEP;
      end select;
    end loop;
  end GAS_STATION_ATTENDANT;
```

Chapter 12

- ***Competition Synchronization with Message Passing:***
 - ***Example--a shared buffer***
 - **Encapsulate the buffer and its operations in a task**
 - **Competition synchronization is implicit in the semantics of `accept` clauses**
 - **Only one `accept` clause in a task can be active at any given time**
- > **SHOW `BUF_TASK` task and the `PRODUCER` and `CONSUMER` tasks that use it (p. 514-515)**

Task Termination

Def: The execution of a task is *completed* if control has reached the end of its code body

- **If a task has created no dependent tasks and is completed, it is terminated**
- **If a task has created dependent tasks and is completed, it is not terminated until all its dependent tasks are terminated**

Chapter 12

- A **terminate clause** in a `select` is just a `terminate statement`
- A **terminate clause** is selected when no `accept clause` is open
- When a `terminate` is selected in a task, the task is terminated only when its master and all of the dependents of its master are either completed or are waiting at a `terminate`
- A block or subprogram is not left until all of its dependent tasks are terminated
- *Priorities*
 - The priority of any task can be set with the `pragma priority`
 - The priority of a task applies to it only when it is in the task ready queue
- *Evaluation of the Ada 83 Tasking Model*
 - If there are no distributed processors with independent memories, monitors and message passing are equally suitable. Otherwise, message passing is clearly superior

Chapter 12

Concurrency in Ada 95

- Ada 95 includes Ada 83 features for concurrency, plus two new features

1. *Protected Objects*

- A more efficient way of implementing shared data
- The idea is to allow access to a shared data structure to be done without rendezvous
- A protected object is similar to an abstract data type
- Access to a protected object is either through messages passed to entries, as with a task, or through protected subprograms
- A *protected procedure* provides mutually exclusive read-write access to protected objects
- A *protected function* provides concurrent read-only access to protected objects
- > **SHOW** the protected buffer code (pp. 518-519)

Chapter 12

2. *Asynchronous Communication*

- **Provided through asynchronous `select` structures**
 - **An asynchronous `select` has two triggering alternatives, and entry clause or a delay**
 - **The entry clause is triggered when sent a message; the delay clause is triggered when its time limit is reached**
- > **SHOW examples (p. 519-520)**

Chapter 12

Java Threads

- The concurrent units in Java are `run` methods
- The `run` method is inherited and overridden in subclasses of the predefined `Thread` class
- *The Thread Class*
 - Includes several methods (besides `run`)
 - `start`, which calls `run` , after which control returns immediately to `start`
 - `yield`, which stops execution of the thread and puts it in the task ready queue
 - `sleep`, which stops execution of the thread and blocks it from execution for the amount of time specified in its parameter
 - `suspend`, which stops execution of the thread until it is restarted with `resume`
 - `resume`, which restarts a suspended thread
 - `stop`, which kills the thread

Chapter 12

- *Competition Synchronization with Java Threads*

- **A method that includes the `synchronized` modifier disallows any other method from running on the object while it is in execution**
- **If only a part of a method must be run without interference, it can be synchronized**

- *Cooperation Synchronization with Java Threads*

- **The `wait` and `notify` methods are defined in `Object`, which is the root class in Java, so all objects inherit them**
- **The `wait` method must be called in a loop**
- ***Example - the queue***
- **---> **SHOW** `Queue` class (p. 524) and the `Producer` and `Consumer` classes (p. 525)**