

# Deadlock

## Comp 305 Lecture 4

## This Week

- ◆ Lecture on transactions and deadlock
- ◆ Tutorial on Lab Exercise 2
- ◆ Homework
  - Chap 6, problems 18 and 19;
  - Chap 7, problems 4, 7 and 9.

## Transactions

- ◆ A critical section ensures mutual exclusion and sequential execution.
- ◆ An atomic transaction gives the effect of a single atomic action.

```
begin transaction
  read X
  if (X >= a) then {
    read Y
    X = X - a
    Y = Y + a
    write X
    write Y
  }
end transaction
```


## Assumptions

- ◆ Three types of memory
  - Volatile
  - Non-volatile
  - Stable
- ◆ Commit
  - All changes become visible
- ◆ Abort
  - Nothing has happened

## Log Based Transactions

```
begin transaction      → <Ti, start>
  read X
  if (X >= a) then {
    read Y
    X = X - a
    Y = Y + a
    write X
    write Y
  }
end transaction        → <Ti, X, oldV, newV>
                       → <Ti, Y, oldV, newV>
                       → <Ti, commit>
```


## How It Works



## Recovery

- ◆ Operations
  - undo(Ti)
  - redo(Ti)
- ◆ Possible states of Ti
  - Committed, but were values written?
  - Aborted, but were values restored?
  - Unfinished
- ◆ Checkpoints

## Recovery At Work



## Concurrent Transactions

<p>T1 begin transaction                  read X                  modify X                  write X                  ...                  do something                  end transaction</p>	<p>←————→</p>	<p>T2 begin transaction                  read X                  read Y                  modify Y using X                  write Y                  end transaction</p>
--	---------------	---

## Conflicting Operations


<p>T1 begin transaction                  read X                  modify X                  write X                  ...                  do something                  end transaction</p>	<p>←————→</p>	<p>T2 begin transaction                  read X                  read Y                  modify Y using X                  write Y                  end transaction</p>
--	---------------	---

abort restores X  
incorrect Y written

## Locks

- ◆ A shared lock allows multiple reads
- ◆ An exclusive lock allows single write
- ◆ Serial schedule
  - All locks are acquired before any are released.

## What is Deadlock?




## Deadlock Defined

- ◆ A set of processes is in a deadlock state when every process in the set is waiting for an event that can only be caused by another process in the set.
- ◆ Necessary Conditions -
  - Mutual exclusion
  - Hold and wait
  - No preemption
  - Circular wait


## Resource Allocation Graph

- ◆ Set of processes [P1, P2, ...]
- ◆ Set of resources [R1, R2, ...]




## Resource Allocation Graph/2

- ◆ Construct a graph by drawing
  - Edges from Pj to Ri on request (Pj, Ri)
  - Edges from Ri to Pj on allocation (Ri, Pj)




## Cycles and Deadlock

- ◆ Deadlock?
  - P1 waiting
  - P2 waiting
  - P3 runnable




## Cycles and Deadlock

- ◆ Deadlock?
  - P1 waiting
  - P2 waiting
  - P3 runnable
- ◆ P3 releases R3
  - P2 may run




## Cycles and Deadlock

- ◆ Deadlock?
  - P1 waiting
  - P2 waiting
  - P3 runnable




## Cycles and Deadlock

- ◆ Deadlock?
  - P1 waiting
  - P2 waiting
  - P3 runnable
- ◆ P3 requests R2




## Cycles and Deadlock

- ◆ Deadlock?
  - P1 waiting
  - P2 waiting
  - P3 runnable
- ◆ P3 requests R2
- ◆ Circular wait
- ◆ Deadlock




## Is Cycle Sufficient?

- ◆ Cycle exists
- ◆ P3 may release R2
- ◆ P1 may continue
- ◆ Cycle broken
- ◆ No deadlock



## Deadlock Prevention

- ◆ Eliminate 1 of 4 necessary conditions.
- ◆ Mutual exclusion
- ◆ Hold and wait
  - request a,b
  - release a,b
  - request c,d
  - release c,d
  - request c,f
  - release c,f
- ◆ No preemption
- ◆ Circular wait




## Deadlock Algorithms

- ◆ Straightforward prevention is inefficient.

	max request	alloc
P0	10	5
P1	4	2
P2	9	2

- ◆ System has 12
- ◆  $\langle P1, P0, P2 \rangle$  is SAFE



## A Formal View

- ◆ A state is *SAFE* iff there exists a sequence of processes  $\langle P0, P1, \dots \rangle$  such that all processes can complete.
- ◆ All other states are *UNSAFE*.
- ◆ All *DEADLOCK* states are *UNSAFE*.
- ◆ From a *SAFE* state system can allocate resource to avoid *DEADLOCK*

## Bankers Algorithm

- ◆ Overview
  - Each process declares maximum need
  - Request granted iff system remains in safe state
- ◆ Definitions
  - Available[1..M]
  - Max[1..N][1..M]
  - Allocation[1..N][1..M]
  - Need[1..N][1..M] = Max - Allocation
  - Request[1..M]

## The Bankers Algorithm

- 1 IF NOT  $Request \leq NEED[i]$  THEN stop
- 2 WHILE  $Available < Request$  DO wait
- 3 Construct a new state by
  - $Available = Available - Request$
  - $Allocation[i] = Allocation[i] + Request$
  - $Need[i] = Need[i] - Request$
- 4 IF new state is unsafe THEN wait

## The Safety Algorithm

- 1  $Work = Available$   
 $Finish = FALSE$
- 2 Find i such that  $P_i$  hasn't finished but could  
 $Finish[i] = FALSE$   
 $Need[i] \leq Work$
- 3 Assume  $P_i$  completes  
 $Work = Work + Allocation[i]$   
 $Finish[i] = TRUE$   
goto 2
- 4 IF for all i  $Finish[i] = TRUE$  THEN SAFE

## Deadlock Detection

- ◆ If we don't avoid it we must detect deadlock.
- ◆ Run Safety Algorithm on current state
  - 2 substitute  
 $Request[i] \leq work$  for  $need[i] \leq Work$
  - 4 If  $Finish[i] = FALSE$  THEN  $P_i$  is deadlocked

## A Deadlock Detection Algorithm

- 1  $Work = Available$   
 $Finish = FALSE$
- 2 Find i such that  $P_i$  hasn't finished but could  
 $Finish[i] = FALSE$   
 $Request[i] \leq Work$
- 3 Assume  $P_i$  completes  
 $Work = Work + Allocation[i]$   
 $Finish[i] = TRUE$   
goto 2
- 4 If  $Finish[i] = FALSE$  THEN  $P_i$  is deadlocked

## Breaking a Deadlock

- ◆ Break circular wait
  - Kill all deadlocked processes
  - Kill deadlocked processes one-by-one
  - Some processes may resist death
- ◆ Preempt resources
  - Implies rollback to a checkpoint