

## This Week

- ◆ Tutorials on lab 3 tomorrow, Thursday.
  - ◆ Homework for 6 May
    - Chap 8, probs 9, 14; Chap 9, prob 10
- A system incorporates a system call, `sleep(t)`, that puts a process to sleep for `t` seconds. Describe how you would implement `sleep()` in Nachos. How accurate would you expect the system call to be, i.e., how close would the difference between the time a process invokes `sleep(t)` and is woken up be to `t`?

## Virtual Memory: Demand Paging


### Comp 305, Lecture 7

© John H. Hine, 1998


## The Size and Nature of Programs

- ◆ Programs
  - Handle many rare exception conditions
  - Allocate more memory than needed
  - Have rarely used features
  - Make distinct passes
- ◆ Advantage in not allocating memory to full program.


## Virtual v Physical Memory




## Demand Paging




## A Page Fault




## A Page Fault



## A Page Fault



## A Page Fault



## Details

- ◆ Which page frames are free?
- ◆ How do I create more free page frames?
- ◆ When are 'writes' reflected on disk?
- ◆ Retrying complex instructions

## Performance of Demand Paging

- ◆ Let  $p$  = probability of a page fault
- ◆ Then effective access time  


$$= (1-p) * \text{memory access} + p * \text{fault access}$$
- ◆ One fault in 200,000 references can cause 10% slow down.
- ◆ Each process uses less memory
- ◆ Other processes use CPU during fault

## Managing Memory

- ◆ Process's memory grows and grows
- ◆ How do we reclaim memory?
  - Process terminates
  - Swap process out
  - Reclaim from active process


## Page Replacement

- ◆ Program with 2 code and 4 data pages




## Page Replacement

- ◆ Program with 2 code and 4 data pages
- ◆ Execute from page 0




## Page Replacement

- ◆ Program with 2 code and 4 data pages
- ◆ Execute from page 0
- ◆ Reference page 2




## Page Replacement

- ◆ Program with 2 code and 4 data pages
- ◆ Execute from page 0
- ◆ Reference page 2
- ◆ ... Reference page 4




## Page Replacement

- ◆ Program with 2 code and 4 data pages
- ◆ Execute from page 0
- ◆ Reference page 2
- ◆ ... Reference page 4
- ◆ ... Reference page 3




## Page Replacement

- ◆ Program with 2 code and 4 data pages
- ◆ Execute from page 0
- ◆ Reference page 2
- ◆ ... Reference page 4
- ◆ ... Reference page 3
- ◆ ...Execute from page 1




## Performance of a Single Process



## Page Replacement Algorithms


- ◆ Within a fixed memory allocation
- ◆ Optimal
  - Replace page used furthest into the future
  - Cannot be implemented
- ◆ FCFS
  - Queue implementation
  - Anomalous behaviour

## FCFS Replacement Anomaly




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient
- ◆ LRU Approximation
  - Set Ordering




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient
- ◆ LRU Approximation
  - Set Ordering




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient
- ◆ LRU Approximation
  - Set Ordering
  - Second Chance




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient
- ◆ LRU Approximation
  - Set Ordering
  - Second Chance




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient
- ◆ LRU Approximation
  - Set Ordering
  - Second Chance




## Replacement Algorithms cont.

- ◆ Least Recently Used
  - Stack implementation
  - Inefficient
- ◆ LRU Approximation
  - Set Ordering
  - Second Chance




## System Performance




## Program Behaviour

- ◆ Do processes always need the same memory?
- ◆ What is optimal performance point?



## Working Set Model

- ◆ Dynamically adjusts frames allocated
- ◆  $W_t(\tau)$  = set of pages referenced in  $(t-\tau, t)$
- ◆  $w_t(\tau) = |W_t(\tau)|$  = working set size at time  $t$



## UNIX Memory Management

- ◆ Clusters = integral number of page frames
- ◆ Core Map
  - free
  - next, prev on free list
  - page type (system, text, data, stack)
  - pointer to owner process
  - location in owner's virtual memory
  - lock bit and lock wanted bit
  - disk location

## Demand Paging

- ◆ Page Allocated from Free List
- ◆ Logical page found and mapped
  - filled from file
  - zero filled
  - paged from swap space
- ◆ Text segments, prefetched by cluster

## Page Replacement

- ◆ Free list is kept full by a clock algorithm
- ◆ Triggered when free list falls below limit
- ◆ Two handed clock

