

This Week


- ◆ Tutorials Thurs: Lab 4 File Systems
- ◆ Homework:
 - Chap 10: problems 1, 7, 10
 - Chap 11: problem 6

I/O and Network Structures

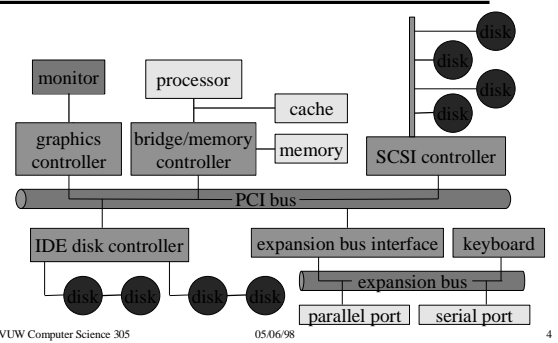
Comp 305 Lecture 10
©John H. Hine 1998

The Problem

- ◆ Extensive use of IO devices, network
- ◆ Devices are idiosyncratic
- ◆ Devices require kernel
- ◆ Speed differences



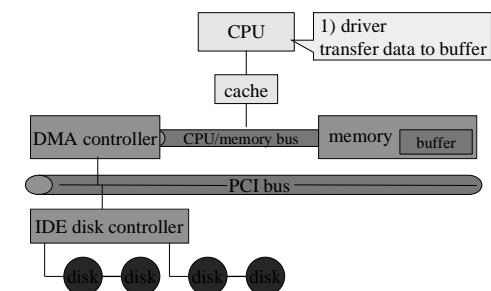
IO Hardware




Controller or Device Interface

- ◆ IO Instructions
- ◆ Memory mapped
- ◆ Polling
- ◆ Interrupts
- ◆ DMA


Direct Memory Access




Direct Memory Access




Direct Memory Access



Direct Memory Access



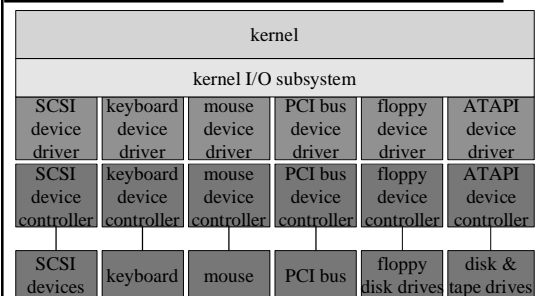
Direct Memory Access



IO Software

- ◆ Device drivers
 - ✦ Block device interface
 - ✦ Character device interface
- ◆ Two level design
 - ✦ Strategy
 - ✦ Detail

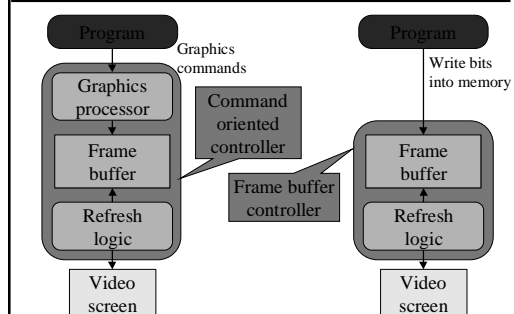
Application Interface



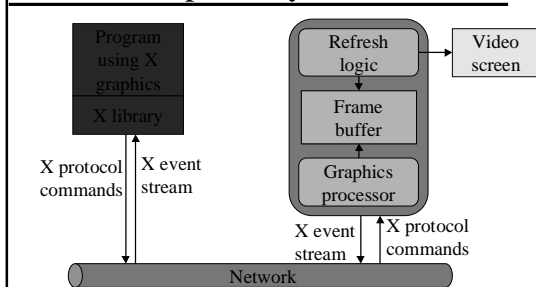
Characteristics

Aspect	Variation	Example

Two Graphics Controller Models




The X Graphics System




IO Subsystem

- ◆ Scheduling
- ◆ Buffering
- ◆ Caching
- ◆ Spooling
- ◆ Error Handling


Double Buffering



Double Buffering




Double Buffering



IO Subsystem

- ◆ Scheduling
- ◆ Buffering
- ◆ Caching
- ◆ Spooling
- ◆ Error Handling

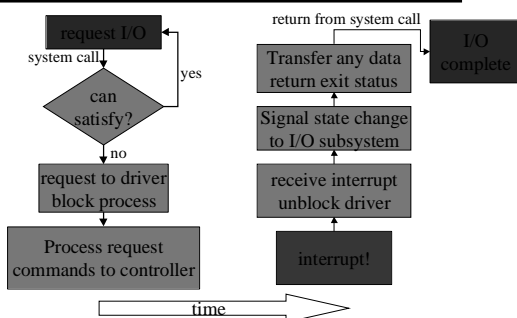
Spooling



IO Subsystem

- ◆ Scheduling
- ◆ Buffering
- ◆ Caching
- ◆ Spooling
- ◆ Error Handling

Making it Happen



Performance Issues

- ◆ Context switches
- ◆ Network IPC
- ◆ Interrupt “readies” a process
- ◆ Encourages big kernel

How to Get Efficient I/O

- ◆ Minimise context switches
 - ✦ Minimise interrupts (polling, large transfers)
- ◆ Minimise data copying
- ◆ Increase concurrency - smart controllers
- ◆ Move processing into IO hardware
- ◆ Balance CPU, memory, bus and IO use

Design Techniques


- ◆ Caching techniques
- ◆ Optimisation and Hints
- ◆ Names

Example Problem: Network File Access

- ◆ Idempotent transactions
- ◆ Remote invocation
 - <read, pathname, offset, length>
- ◆ Requires
 - ✦ Multiple directory lookups
 - ✦ Read inode or similar
 - ✦ Read data block

Caching

- ◆ Remember previous result of operation
 - ✦ Virtual memory
 - ✦ TLB
 - ✦ Path name look up
 - ✦ File descriptor table
 - ✦ Disk block caches



Issues With Caching

- ◆ Searching for cached results
- ◆ Cache invalidation
 - ✦ Hooks
- ◆ Cache replacement
- ◆ Locality of operation

Optimisation and Hints


- ◆ Optimisation depends on predicting the future
- ◆ Approaches to optimisation:
 - ✦ Every instance
 - » Fixing a linear search by switching to a binary search (always faster)
 - ✦ Optimise some instances
 - » Caching results (sometimes faster)

Hints

- ◆ Hint: Answer that may be correct, but isn't guaranteed
- ◆ Examples:
 - ✦ Name address binding
 - ✦ Resource location
- ◆ Difference from caching
 - ✦ Try hint to determine correctness


Hierarchical Names

- ◆ Name = identifier
- ◆ Object - to be identified
- ◆ Map: Names onto objects



Name Hierarchy

- ◆ Examples
 - ✦ Domain names
 - ✦ IP addresses
 - ✦ Files
 - ✦ X.500



Types of Names

- ◆ External
 - ✦ Character strings
 - ✦ User friendly
- ◆ Internal
 - ✦ Bit string
 - ✦ May contain hint to location
- ◆ Unique