

UNIX Design Principles

- Designed to be a time-sharing system.
- Has a simple standard user interface (shell) that can be replaced.
- File system with multilevel tree-structured directories.
- Files are supported by the kernel as unstructured sequences of bytes.
- Supports multiple processes; a process can easily create new processes.
- High priority given to making system interactive, and providing facilities for program development.

Programmer Interface

Like most computer systems, UNIX consists of two separable parts:

- Kernel: everything below the system-call interface and above the physical hardware.
 - Provides file system, CPU scheduling, memory management, and other OS functions through system calls.
- Systems programs: use the kernel-supported system calls to provide useful functions, such as compilation and file manipulation.

4.3BSD Layer Structure

(the users)		
shells and commands compilers and interpreters system libraries		
<i>system-call interface to the kernel</i>		
signals terminal handling character I/O system terminal drivers	file system swapping block I/O system disk and tape drivers	CPU scheduling page replacement demand paging virtual memory
<i>kernal interface to the hardware</i>		
terminal controllers terminals	device controllers disks and tapes	memory controllers physical memory

Process Control Blocks

- The most basic data structure associated with processes is the *process structure*.
 - unique process identifier
 - scheduling information (e.g., priority)
 - pointers to other control blocks
- The *virtual address space* of a user process is divided into text (program code), data, and stack segments.
- Every process with sharable text has a pointer from its process structure to a *text structure*.
 - always resident in main memory
 - records how many processes are using the text segment
 - records where the page table for the text segment can be found on disk when it is swapped

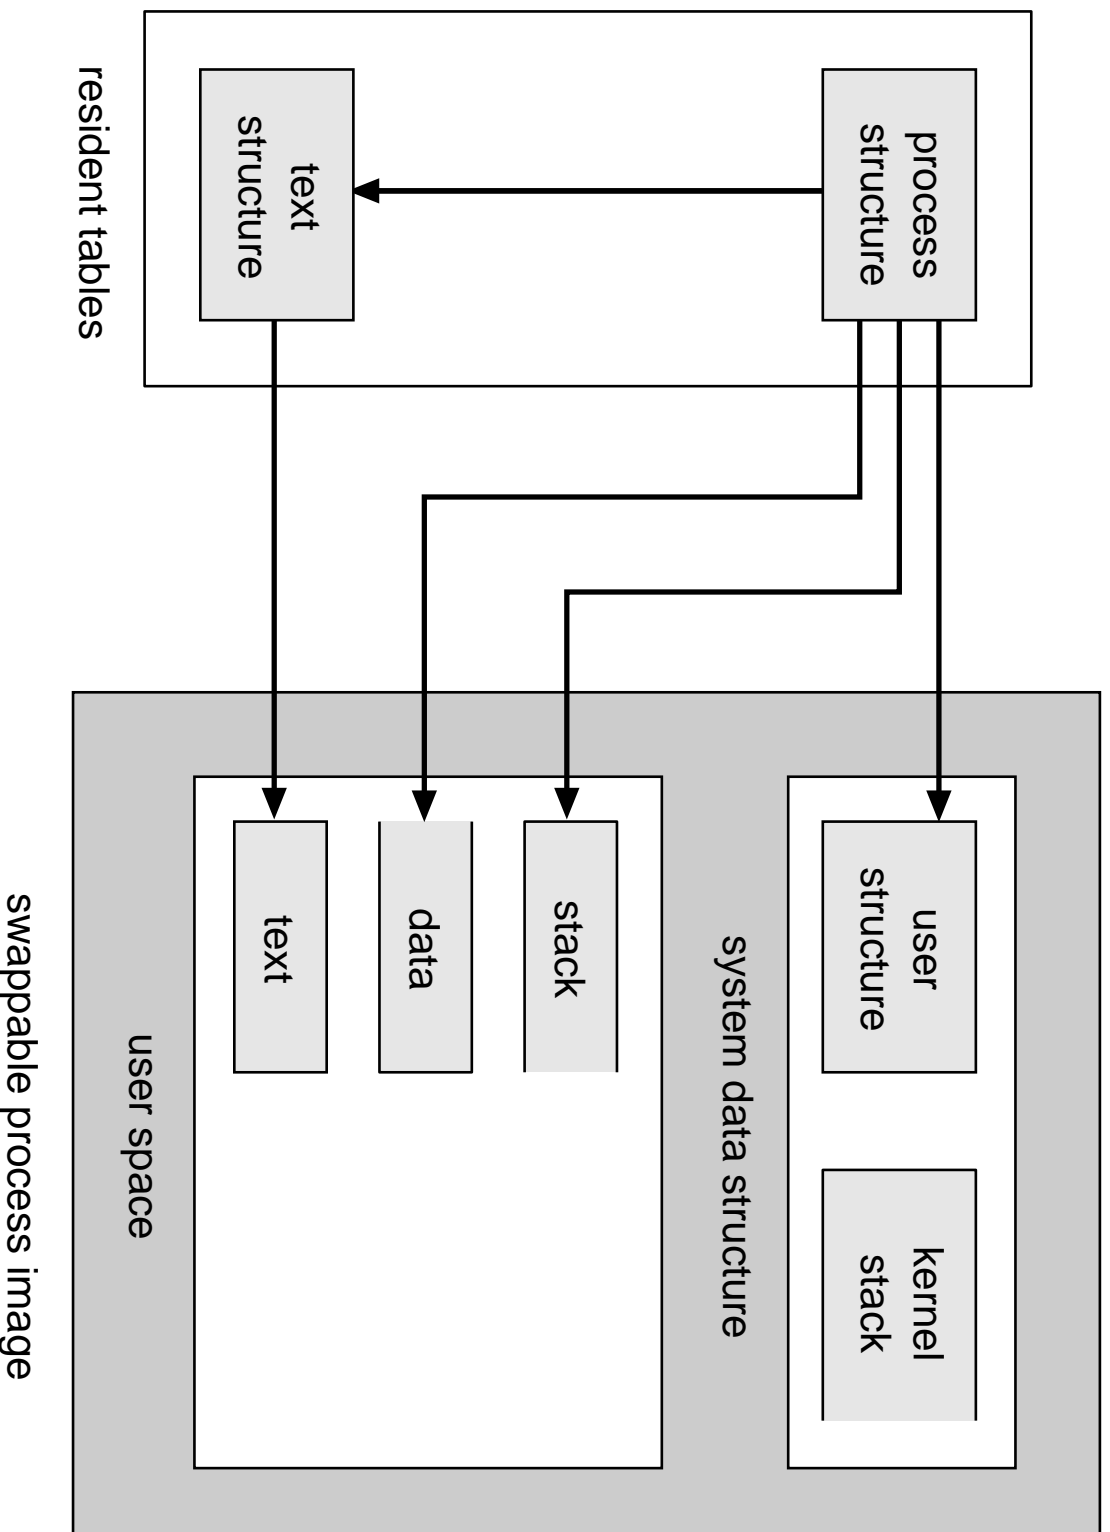
Process Control Blocks (Cont'd)

- The *page tables* record information on the mapping from the process' virtual memory to physical memory.
- Information about the process that is needed only when the process is resident is kept in the *user structure* (or *u structure*).
 - mapped read-only into user virtual address space
 - writable by the kernel
 - maintains the current directory and the table of open files

System Data Segment

- Most ordinary work is done in *user mode*; system calls are performed in *system mode*.
- The system and user phases of a process never execute simultaneously.
- A *kernel stack* (rather than the user stack) is used for a process executing in system mode.
- The kernel stack and the user structure together compose the *system data segment* for the process.

Finding parts of a process using process structure



Allocating a New Process Structure

- **fork** allocates a new process structure for the child process, and copies the user structure.
 - new page table is constructed
 - new main memory is allocated for the data and stack segments of the child process
 - copying the user structure preserves open file descriptors, user and group identifiers, signal handling, etc.

Allocating a New Process Structure (Cont'd)

- **fork** does *not* copy the data and stack to the new process; the new process simply shares the page table of the old one.
 - new user structure and a new process structure are still created
 - commonly used by a shell to execute a command and to wait for its completion
- A parent process uses **fork** to produce a child process; the child uses **execve** to change its virtual address space, so there is no need for a copy of the parent.
- Using **fork** with a large parent process saves CPU time, but can be dangerous since any memory change occurs in both processes until **execve** occurs.
- **execve** creates no new process or user structure; rather, the text and data of the process are replaced.

Memory Management

- The initial memory management schemes were constrained in size by the relatively small memory resources of the PDP machines on which UNIX was developed.
- Pre-3BSD systems use swapping exclusively to handle memory contention among processes: If there is too much contention, processes are swapped out until enough memory is available.
- Allocation of both main memory and swap space is done first-fit.

Memory Management (Cont'd)

- Sharable text segments do not need to be swapped; results in less swap traffic and reduces the amount of main memory required for multiple processes using the same text segment.
- The *scheduler process* (or *swapper*) decides which processes to swap in or out, considering such factors as time idle, time in or out of main memory, size, etc.
- In 4.3BSD, swap space is allocated in pieces that are multiples of a power of 2 and a minimum size, up to a maximum size determined by the size of the swap-space partition on the disk.

Paging

- Berkeley UNIX systems depend primarily on paging for memory-contention management, and depend only secondarily on swapping.
- *Demand paging* – When a process needs a page and the page is not there, a page fault to the kernel occurs, a frame of main memory is allocated, and the proper disk page is read into the frame.
- A *pagedaemon* process uses a modified second-chance page-replacement algorithm to keep enough free frames to support the executing processes.
- If the *scheduler* decides that the paging system is overloaded, processes will be swapped out whole until the overload is relieved.