

Module 11

Objects and Relations

© Louis D. Nel 1996

95.305

Objects and Relations 11 - 1

95.305

Objectives

- **Learn about issues relating to objects and relations**

© Louis D. Nel 1996

95.305

Objects and Relations 11 - 2

Topics

- **Programs and Databases: Persistence**
- Proxies and Adapters
- Storing Objects in Relational Tables
- Performance Issues


References

- **White paper on TOPLINK from ObjectPeople**
<http://www.objectpeople.com>
- **Ivar Jacobson, Object-Oriented Software Engineering: A use case driven approach, Addison Wesley, 1992, Chapter 10**

Topics

- Programs and Databases: Persistence
- **Proxies and Adapters**
- Storing Objects in Relational Tables
- Performance Issues

Making Objects Persistent




- **Most large applications have persistent data**
- **An object-oriented application wants to maintain persistent objects**
- **Using a database to store persistent objects has advantages**
 - database provide integrity
 - support multi-users, concurrency
 - may be existing legacy system

- **Storing persistent objects in Object-Oriented Database has many advantages**
 - database stores objects (not tables)
 - translation between persistent entities and local entities is easy (both are objects)
 - good match between programming language features and data model features (inheritance, data and behaviour, ...)
- **But, OODBMS's are still in their infancy and not well established like relational databases**
- **OODBMS's can be slow**

Storing Objects in a Relational Database


- Storing Objects in a relational database causes some problems
 - stores tables not objects
 - “impedance mismatch” between application data model and that of database
- Still many people are trying to do just that

Approach 1



- Objects must know about relational structure and database querying with SQL
- Application objects get polluted with database issues

Approach 2



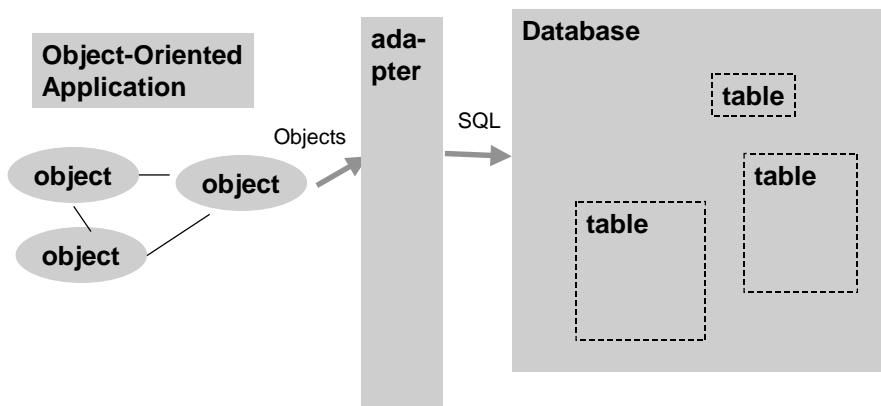
- Pass objects to database to smash into tables
- But, RDBM's don't understand objects; not an industry standard interface like SQL

© Louis D. Nel 1996

95.305

Objects and Relations 11 - 11

Approach 3



- Use dedicated adaptor between objects and tables
- At least application does not have to know about tables, and RDBMS does not know about objects

© Louis D. Nel 1996

95.305

Objects and Relations 11 - 12

TOPLink

The diagram illustrates the TOPLink architecture. On the left, a box labeled "Smalltalk Program" contains three oval shapes, each labeled "object". Lines connect these objects, and an arrow labeled "Objects" points from them to a central vertical bar labeled "TOP-Link". From the "TOP-Link" bar, an arrow labeled "SQL" points to a large box on the right labeled "Database". Inside the "Database" box, there are three dashed-line rectangles, each labeled "table".

- **TOPLink, from ObjectPeople, is such an adapter for Smalltalk programs**

© Louis D. Nel 199695.305Objects and Relations 11 - 13

95.305

Topics

- **Programs and Databases: Persistence**
- **Proxies and Adapters**
- **Storing Objects in Relational Tables**
- **Performance Issues**

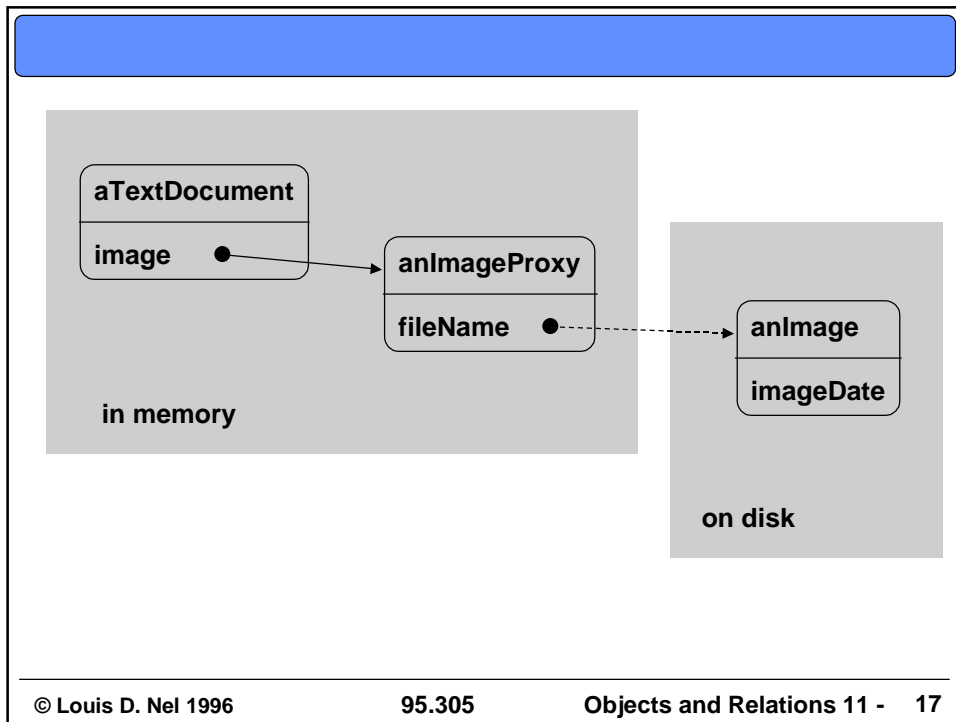
© Louis D. Nel 199695.305Objects and Relations 11 - 14

Proxy Pattern

- **Intent:**
Provide a surrogate or placeholder for another object to control access to it, to introduce side effects for messages sent to the subject.
- **Aka:**
Surrogate, Encapsulator
- **Other Structural Patterns**
Adapter, Bridge, Composite, Decorator, Facade, Flyweight

Proxy Pattern -Motivation

- **Example: Deferred creation or initialization**
- Text document with graphical objects embedded may not want to create the graphical object all when the document is opened, perhaps only those which are in view.
- Suggests using a less expensive proxy in place of the real image so the document editor does not know the difference.
- The proxy behaves just like the real image, but only creates it when it needs to be drawn (e.g. when draw method is invoked).



Proxy Pattern: Participants and Collaboration

- **Proxy**
 - maintains reference to real subject
 - provides an interface identical to subjects (client does not know its dealing with a proxy)
 - control access to real subject (and may have to create or destroy it)
- **Proxy forwards request to real subject**
- **Remote proxy maps & encodes request into new address space**
- **Virtual proxy may cache info about real subject to postpone access**
- **Protection proxy ensures client has right to access real subject**

Proxy Pattern: Copy-on-Write example


- Copying large objects can be expensive
- Idea: only copy an object if the copy gets modified
- Client thinks they have a copy, but they actually have a proxy to the original
- Only copy the original if client actually tries to modify the proxy (perform a write operation)

Proxy Pattern: Related Patterns

- Adapter
Provides a different interface to an object.
- Adapts (changes) the interface to an object, proxies offer the same interface
- (Protection proxies though may refuse a certain interface request)
- Decorator
Similar to proxies but add one or more responsibilities to that of subject
- So, decorators have a different purpose

Related Previous Work: Encapsulators

- Pascoe, Geoffrey A. "Encapsulators: A New Software Paradigm in Smalltalk-80, OOPSLA '86



Adapter Transparency

- Adapter should be transparent:
- Application objects should not be forced to be a subclass of `Adapter`
- Could be implemented with transparent Proxy techniques

Topics

- **Programs and Databases: Persistence**
- **Proxies and Adapters**
- **Storing Objects in Relational Tables**
- **Performance Issues**

Issues an Object-RDBMS Adapter must address

- **Object Identity**
- **In an OOP objects have identity irrespective of their attribute values -no two objects are identical**
- **Relational models handles uniqueness through primary keys**
- **Adapter must match the key in the table with an objects internal identity when in local memory**

Storing Objects as Data in Tables

Customer
name = 'Lou'
account = 1001
(id = 101)

Customer
name = 'Sue'
account = 1231
(id = 112)

CUSTOMER		
<u>ID</u>	NAME	ACCOUNT
101	Lou	1001
112	Sue	1231

- Simple idea: Create table for class and use internal (object) id for key
- Note: internal object id is not visible to the OO application programmer

Nested Objects

Customer
name = 'Lou'
account
(id = 101)

BankAccount
accountNumber=1001
balance = 900
owner = 'Lou'
(id = 100)

CUSTOMER		
<u>ID</u>	NAME	ACCOUNT
101	Lou'	100
112	Sue'	111


Customer
name = 'Sue'
account
(id = 112)

BankAccount
accountNumber= 1231
balance = 950
owner = 'Sue'
(id = 111)

ACCOUNT			
<u>ID</u>	ACCTNUM	BALANCE	OWNER
100	1001	900	Lou'
111	1231	950	Mary'

- Nested objects stored in separate tables with foreign keys

OO Application Code Example



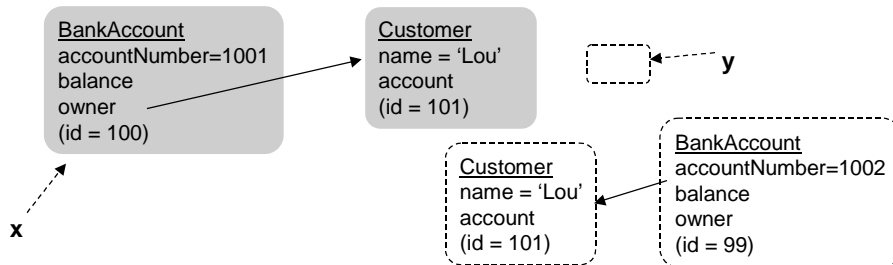
```

| cust acct |
cust := Customer new name: 'Lou'.
x := BankAccount new owner: cust.
y := BankAccount new owner: cust.
x customer == y customer
true
  
```

© Louis D. Nel 1996

95.305

Objects and Relations 11 - 27



```

| cust acct |
cust := Customer new name: 'Lou'.
x := BankAccount new owner: cust.
y := BankAccount new owner: cust.
y storeOnDatabase
  
```

© Louis D. Nel 1996

95.305

Objects and Relations 11 - 28

"Object instance method"

storeOnDatabase

"make receiver persistent and remove receiver"

```
| key |  
key := Database insert: self asTuple.  
self become:  
    (Proxy new class: self class;  
        key: key).
```

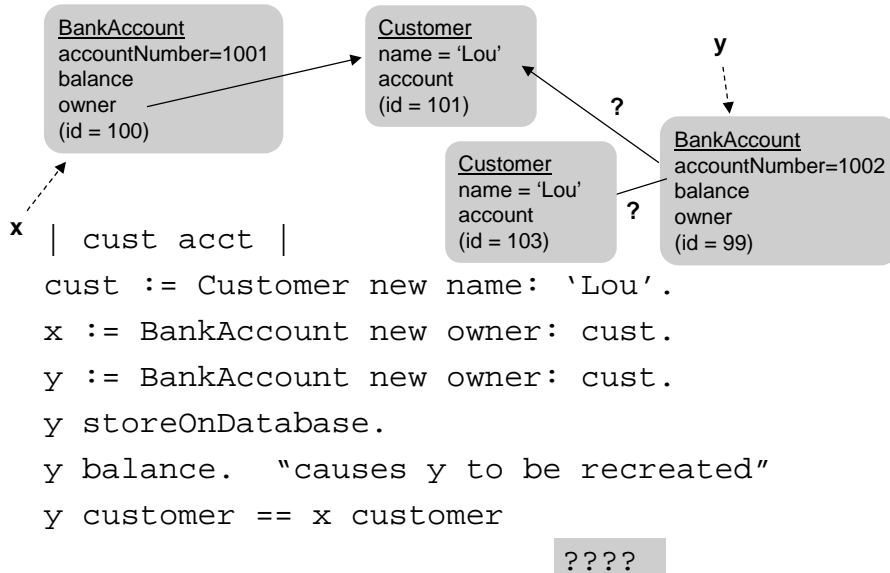
"Proxy instance method"

doesNotUnderstand

"replace receiver proxy with object created from database"

```
|object tuple|  
object := self subject class new.  
tuple := Database select: 'all'  
                from: subject class name  
                where: id = subject key.  
1 to tuple size do:  
    [:i|object instVar at: i put: tuple at: i]  
self become: object
```

Are object references recovered properly?

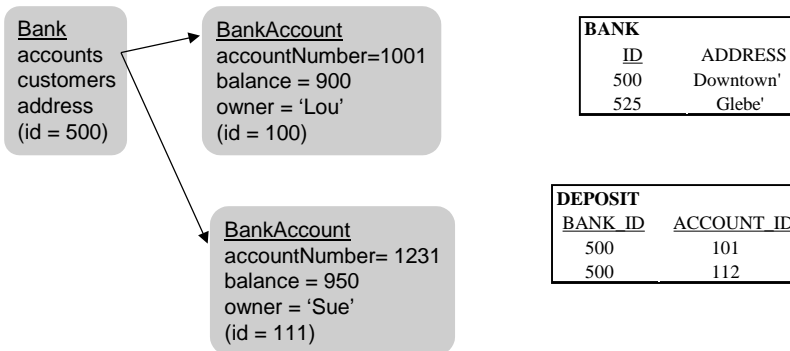


© Louis D. Nel 1996

95.305

Objects and Relations 11 - 31

Collections (1:N Relationships)



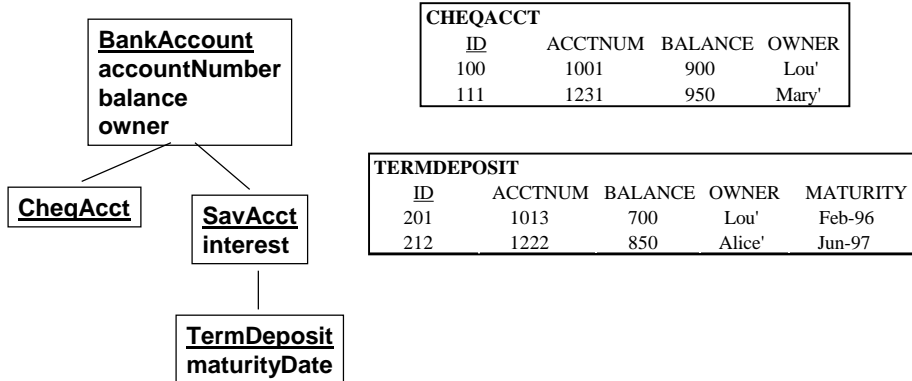
- **1:N Relationships (Collection Attributes) might be stored in separate tables**

© Louis D. Nel 1996

95.305

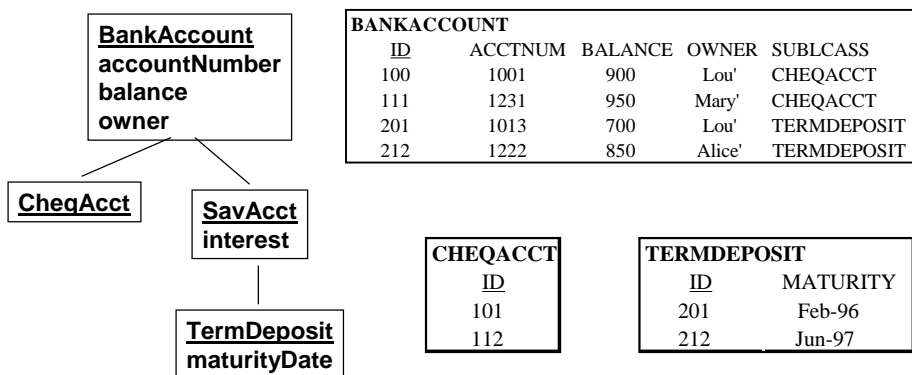
Objects and Relations 11 - 32

What about Inheritance



- Option1: Represent concrete classes with inherited attributes, but don't represent abstract classes

...What about Inheritance



- Option2: Store common attributes in Abstract class table and inherited attributes in subclass tables

Normalization

- Jacobson suggests:
- A Database designed based on objects will normally end up in 3NF
- “Because table represent objects with unique id and mutually exclusive facts about that object”
- “Because OO models often model reality we tend to identify unique objects and assign attributes where they belong”
- may be a bit naive

95.305

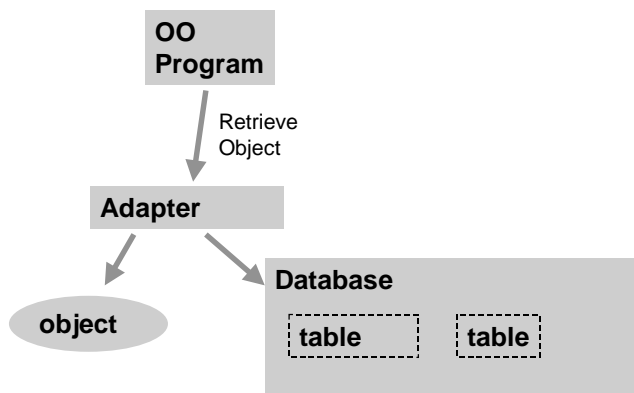
Topics

- Programs and Databases: Persistence
- Proxies and Adapters
- Storing Objects in Relational Tables
- Performance Issues

Adapter Performance Issues

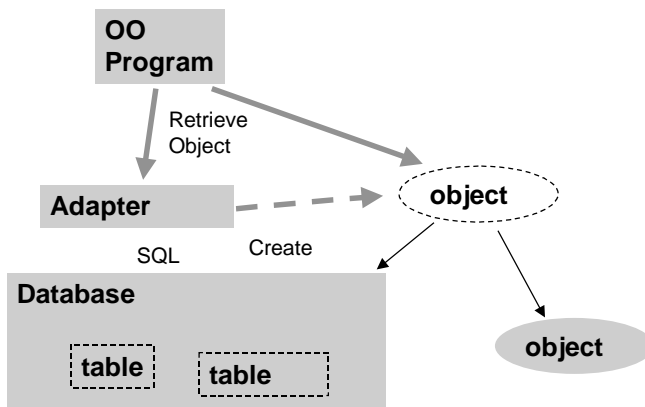
- A Query to the RDBMS might result in many objects.
- The adapter should probably retrieve only a few and provide proxies for others until they are actually needed
- Object proxies can be constructed by transparent wrapper objects
- Adapter should also Cache objects so subsequent references don't go to the database (The application-database interactions are the expensive thing)

Caching



- If the adapter is asked to retrieve an object it first checks to see if it is already in the cache

Caching



- Adapter creates a proxy which the OO program references
- Proxies may have cached objects or else talk to database

Large Results

- Query to database can result in a large number of objects
- Building objects for all rows in a resulting table can be too expensive
- Adapter can use a Stream with cursor position to scroll through “objects” and only create them when needed (amortise over the their use)
- Reduces the perceived delay of fetching objects from the database

Ad hoc Queries from within Objects

- Occasionally it might be desirable for Objects to know about the database and make direct SQL queries
- TOPLink, for example, provides a SQL SELECT-like access using familiar Smalltalk blocks
- e.g.

```
aTopLinkSession
  readAllFor: Employee
  where: [:anEmployee |
    anEmployee manager address street = 'Elm'
    and: (anEmployee lastName = 'Smith') ]
```

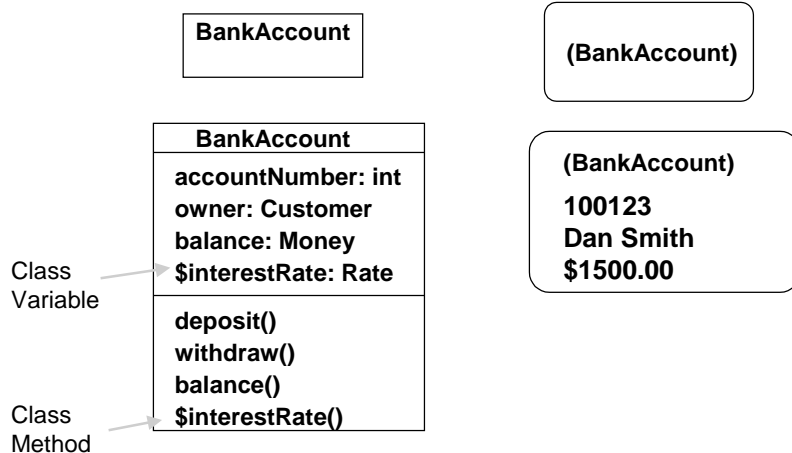
Reading and Writing Data

- It is one thing to keep your constructed objects consistent with the data in the database,
- It is much more difficult to ensure when writing data that database stays consistent with object model

Summary

- The issues here have been greatly over simplified
- It is a very challenging problem to store objects in relational databases and recover them properly
- Especially tricky with recursive references

Showing Class and Instance Details



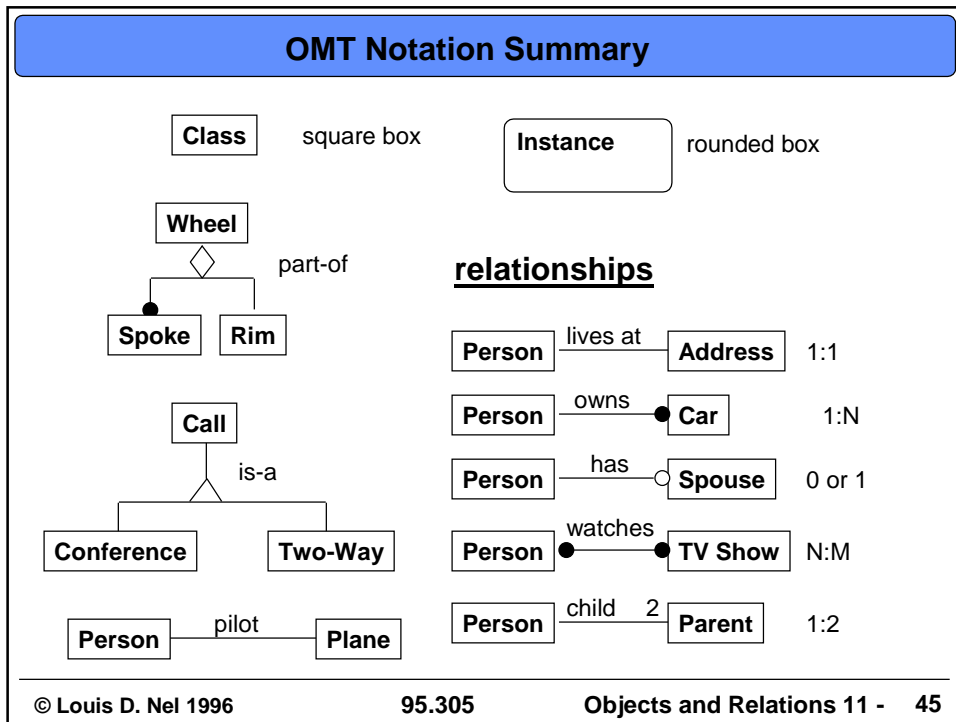
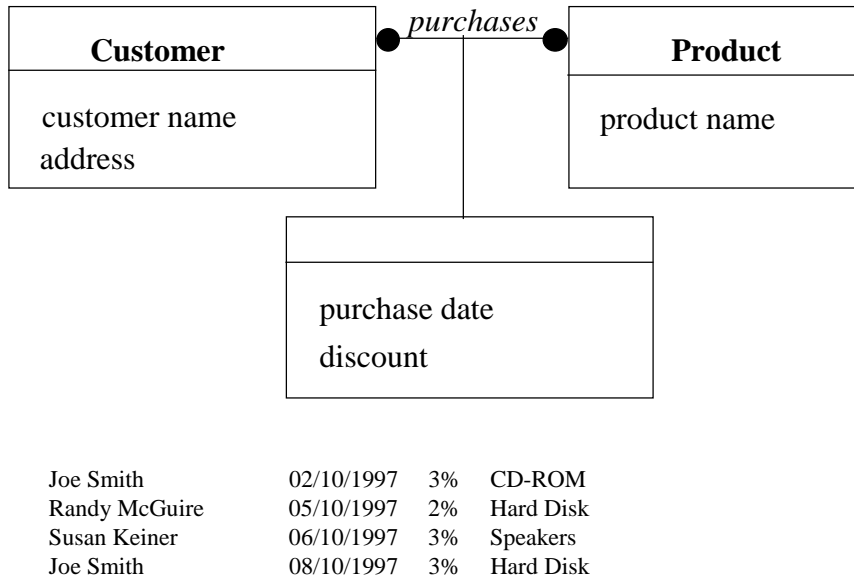


Figure 12. Mapping Classes to Tables			
Object Model	<div>Patient patient name sin address</div>		
Table Model	Attribute name	Nulls?	Domain
	patient-ID	N	long
	patient name	N	string
	sin	Y	long
	address	Y	string
SQL Code	CREATE TABLE Patient		
	(patient-ID long not null, patient -name char(30) not null, social-insurance-number long, address char(30), PRIMARY KEY (patient-ID));		
© Louis D. Nel 1996			
95.305			
Objects and Relations 11 - 46			

Figure 5. Link Attributes

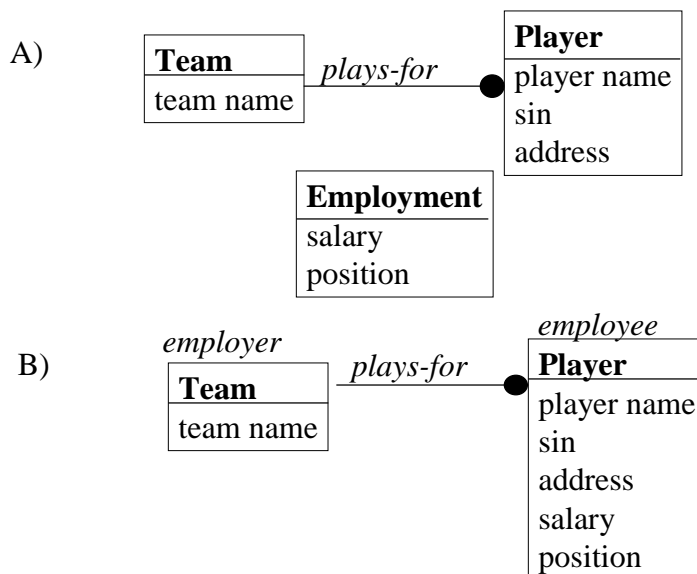


© Louis D. Nel 1996

95.305

Objects and Relations 11 - 47

Figure 6. Link Versus Object Attributes

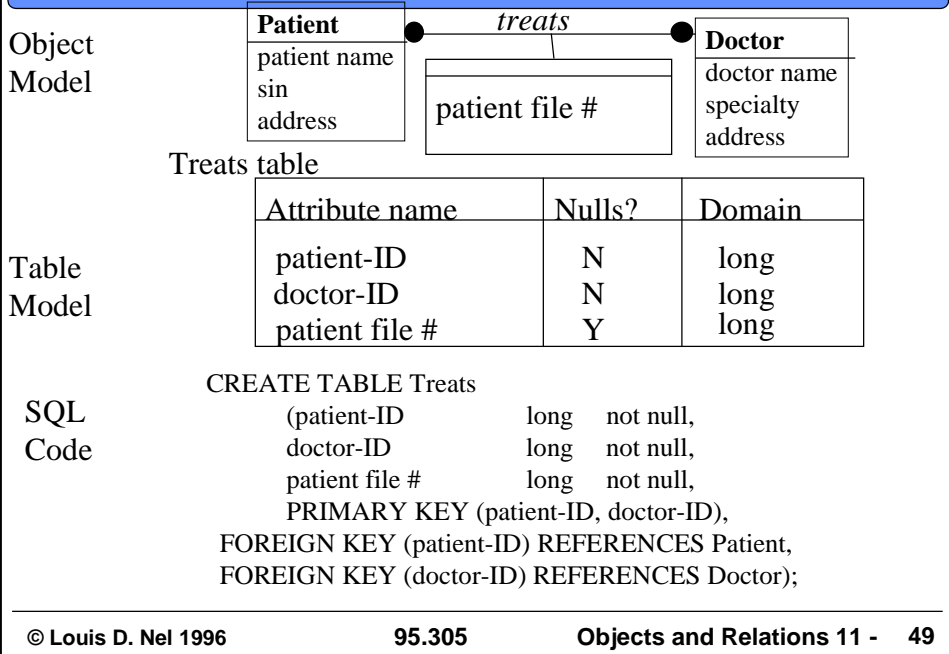


© Louis D. Nel 1996

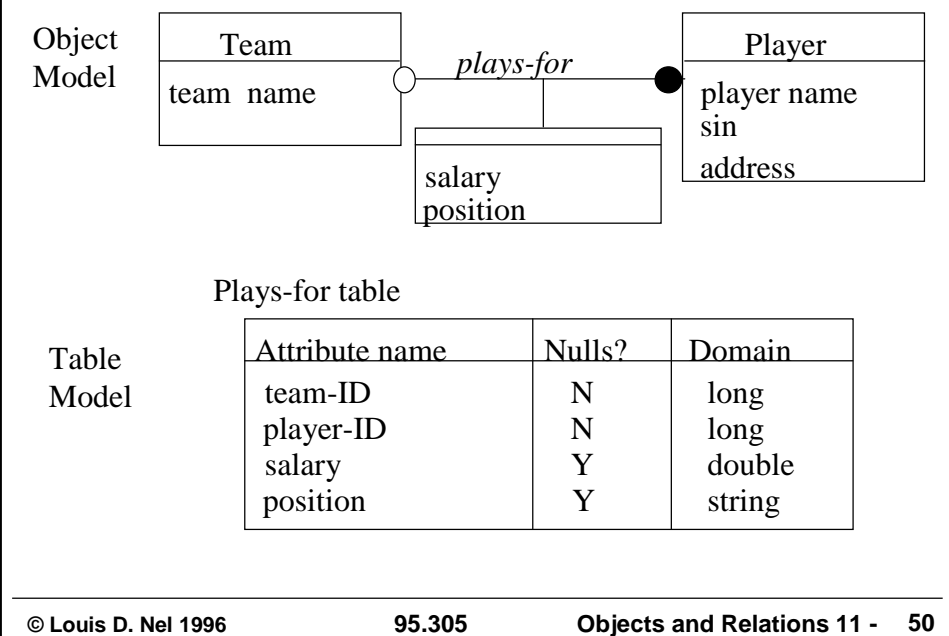
95.305

Objects and Relations 11 - 48

Figure 13. Mapping Many-To-Many Associations to Tables



Mapping 1:N Associations to Tables, distinct association table



Mapping 1:N Associations to Tables, buried foreign key

Table Model

Player table

Attribute name	Nulls?	Domain
player-ID	N	long
player name	N	string
sin	Y	long
address	Y	string
team-ID	Y	long
salary	Y	double
position	Y	string

Figure 9. Investment Generalization Relationship

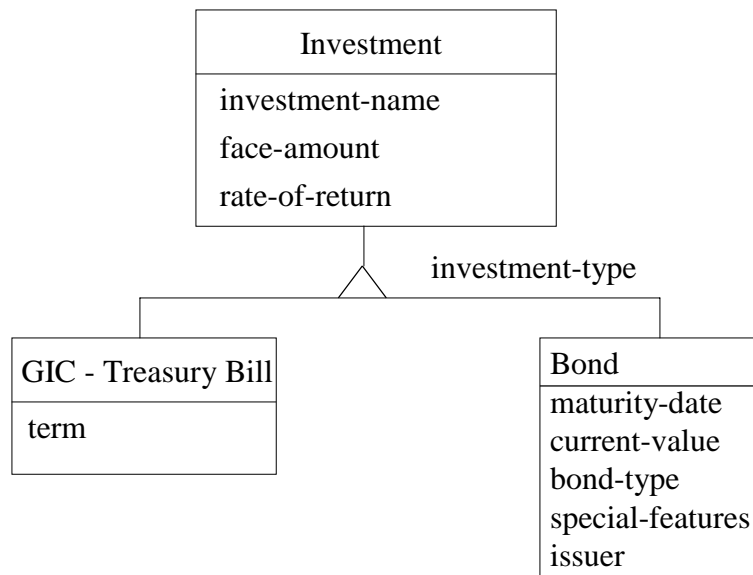


Figure 19. Mapping Generalizations to Tables, superclass & subclass tables			
Investment Table	Attribute name	Nulls?	Domain
	investment-ID	N	long
	investment-name	N	string
	face-amount	Y	double
	rate-of-return	Y	double
	investment-type	N	string
Bond Table	Attribute name	Nulls?	Domain
	investment-ID	N	long
	maturity-date	Y	date
	current-value	Y	double
	bond-type	Y	string
	special-features	Y	string
	issuer	Y	string
GIC - Treasury Bill Table	Attribute name	Nulls?	Domain
	investment-ID	N	long
	term	Y	long
© Louis D. Nel 1996 95.305 Objects and Relations 11 - 53			

Figure 20. Mapping Generalizations to Tables, many subclass tables			
GIC - Treasury Bill Table	Attribute name	Nulls?	Domain
	investment-ID	N	long
	investment-name	N	string
	face-amount	Y	double
	rate-of-return	Y	double
	term	Y	long
Bond Table	Attribute name	Nulls?	Domain
	investment-ID	N	long
	investment-name	N	string
	face-amount	Y	double
	rate-of-return	Y	double
	maturity-date	Y	date
	current-value	Y	double
	bond-type	Y	string
	special-features	Y	string
	issuer	Y	string
© Louis D. Nel 1996 95.305 Objects and Relations 11 - 54			

Figure 21. Mapping Generalizations to Tables, one superclass table

Investment Table	Attribute name	Nulls?	Domain
	investment-ID	N	long
	investment-name	N	string
	face-amount	Y	double
	rate-of-return	Y	double
	investment-type	N	string
	term	Y	long
	maturity-date	Y	date
	current-value	Y	double
	bond-type	Y	string
	special-features	Y	string
	issuer	Y	string

Figure 8. Aggregation

