

# Module 12

## Query Processing and Optimization

© Louis D. Nel 1996

95.305

Query Processing 12 -

1

95.305

### Objectives

- Learn about issues affecting how queries should be processed

© Louis D. Nel 1996

95.305

Query Processing 12 -

2

**95.305**

## **Topics**

## **References**

- **Elmasri & Navathe Chapter 16**

## Query Processing Schematic

**Query (high level language -SQL)**

Scanning  
Parsing and  
Validating

validating ensure tables, attributes exist in database

**Intermediate query form (query tree, query graph)**

Query Optimizer

formulate an efficient execution strategy

**Execution Plan**

Query code generator

**Code to execute Query**

run-time Database Processor

interpreted or compiled

**Query Result**

© Louis D. Nel 1996

95.305

Query Processing 12 -

5

## Optimizing Declarative Queries

- **Procedural languages are optimized by the query programmer since the instruct how the database is to be navigated**
- **Declarative languages indicate what data is required, and not how to navigate the database (e.g. SQL)**
- **For declarative languages an optimization strategy can greatly improve the run-time efficiency**

© Louis D. Nel 1996

95.305

Query Processing 12 -

6

## Optimization

- The optimizer has several implementations for operations like SELECT, JOIN to pick from
- Two main optimization techniques: Heuristics and Cost estimation
- Heuristics typically reorder queries in the execution plan (query tree)
- Estimation is based on computing a probable cost of various strategies
- The two techniques are usually combined by the query optimizer

## Query Operations algorithms

- The optimizer has several algorithms for implementing SELECT, JOIN, etc.
- Choice of algorithm might depend on:
  - size of tables
  - knowledge of data in tables
  - whether an index for the appropriate attribute exists
  - whether a hash function exists for the query attributes
  - ...

## Company Relational Database Schema

### EMPLOYEE

fname	minit	lname	<u>ssn</u>	bdate	address	sex	salary	superssn	dno
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

### DEPARTMENT

dname	<u>dnumber</u>	mgrssn	mgrstartdate
-------	----------------	--------	--------------

### DEPT\_LOCATIONS

<u>dnumber</u>	<u>dlocation</u>
----------------	------------------

### PROJECT

pname	<u>pnumber</u>	plocation	dnum
-------	----------------	-----------	------

### WORKS\_ON

<u>essn</u>	pno	hours
-------------	-----	-------

### DEPENDENT

<u>essn</u>	<u>dependent_name</u>	sex	bdate	relationship
-------------	-----------------------	-----	-------	--------------

**Note:**

Attributes referring to the same thing can have different name  
(pnumber vs. pno)

Attributes referring to different things can have the same name  
( name of employee or name of department )

## Company Relations

### EMPLOYEE

FNAME	INIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALAR	<u>SUPERSSN</u>	DNO
John	B	Smith	123456789	9-Jan-55	731 Fondern	M	30000	333445555	5
Franklin	T	Wong	333445555	8-Dec-45	638 Voss	M	40000	888665555	5
Alicia	J	Zelaya	999887777	19-Jul-58	3321 Castle	F	25000	987987987	4
Jennifer	S	Wallace	987654321	20-Jun-31	291 Berry	F	43000	888665555	4
Ramesh	K	Narayan	666884444	15-Sep-52	975 Fire Oak	M	38000	333445555	5
Joyce	A	English	453453453	31-Jul-62	5631 Rice	F	25000	333445555	5
Ahmad	V	Jabber	987987987	29-Mar-59	980 Dallas	M	25000	987654321	4
James	E	Borg	888665555	10-Nov-27	450 Stone	M	55000	NULL	1

### DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
Research	5	333445555	22-May-78
Administration	4	987654321	1-Jan-85
Headquarters	1	888665555	19-Jun-71

## ...Company Relations

**DEPT\_LOCATIONS**

<u>DNUMBER</u>	<u>DLOCATION</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

**WORKS\_ON**

<u>ESSN</u>	<u>PNO</u>	<u>HOURS</u>
123456789	1	32.50
123456789	2	7.50
666884444	3	40.00
453453453	1	20.00
453453453	2	20.00
333445555	2	10.00
333445555	3	10.00
333445555	10	10.00
333445555	20	10.00
999887777	30	30.00
999887777	10	10.00
987987987	10	35.00
987987987	30	5.00
987654321	30	20.00
987654321	20	15.00
888665555	20	NULL

© Louis D. Nel 1996

95.305

Query Processing 12 -

11

## ...Company Relations

**PROJECT**

<u>PNAME</u>	<u>PNUMBER</u>	<u>PLOCATION</u>	<u>DNUM</u>
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
NewBenefits	30	Stafford	4

**DEPENDENT**

<u>ESSN</u>	<u>D_NAME</u>	<u>SEX</u>	<u>BDATE</u>	<u>RELATIONSHIP</u>
333445555	Alice	F	5-Apr-76	DAUGHTER
333445555	Theodore	M	25-Oct-73	SON
333445555	Joy	F	3-May-48	SPOUSE
987654321	Abner	M	29-Feb-32	SPOUSE
123456789	Michael	M	1-Jan-78	SON
123456789	Alice	F	31-Jan-78	DAUGHTER
123456789	Elizabeth	F	5-May-57	SPOUSE

© Louis D. Nel 1996

95.305

Query Processing 12 -

12

## Implementations for SELECT

- We will consider several select operations on the company database

$\sigma_{(ssn=123456789)}$  (EMPLOYEE)

$\sigma_{(DNUMBER > 5)}$  (DEPARTMENT)

$\sigma_{(DNO = 5)}$  (EMPLOYEE)

$\sigma_{(DNO = 5 \text{ AND } SALARY > 30000 \text{ AND } SEX = F)}$  (EMPLOYEE)

$\sigma_{(ESSN=123456789 \text{ AND } PNO=10)}$  (WORKS\_ON)

## Implementation of SELECT search

- Brute force (Linear Search) Retrieve every record from the file and test whether its attribute values satisfy the select condition
- Works for any SELECT query
- Makes no assumptions about file structure
- Obviously quite slow:  $O(n)$  for  $n$  records

### ...Implementation of SELECT search

$\sigma_{(ssn=123456789)}$  (EMPLOYEE)

- **Binary Search**
- **If the selection condition involves equality on the ordering key binary search can be used**
- **$O(\log(n))$  for n records**

### ...Implementation of SELECT search

$\sigma_{(ssn=123456789)}$  (EMPLOYEE)

- **Using primary key or hash function**
- **Works is selection condition is equality test on search key**
- **usually faster than  $O(\log(n))$  for n records**
- **Could be as fast as  $O(1)$**



### ...Implementation of SELECT search

$\sigma_{(DNUMBER > 5)}(DEPARTMENT)$

- Use primary index to retrieve multiple records
- Works if selection condition is an inequality on search key
- Use the index to find the corresponding equality condition ( $DNUMBER = 5$ ) then chain through subsequent records
- Can only be used for range queries (e.g.  $30000 \leq SALARY \leq 35000$ )

### ...Implementation of SELECT search

$\sigma_{(DNO = 5)}(EMPLOYEE)$

- Use a clustering index to retrieve multiple records
- Works for selection conditions involving equality on a nonkey attribute with a clustering index

### ...Implementation of SELECT search

$\sigma_{(DNO = 5)}(EMPLOYEE)$

- Use a secondary (B+ tree) index
- Works for retrieving a single record if indexing field has unique value (is key) or to retrieve multiple records if indexing field is not a key
- Can only be used for range queries (e.g.  $30000 \leq SALARY \leq 35000$ )

### ...Implementation of SELECT search

$\sigma_{(DNO = 5 \text{ AND } SALARY > 30000 \text{ AND } SEX = F)}(EMPLOYEE)$

- Conjunctive Selection
- If the condition is of the form  $X \text{ AND } Y$ , use one for the previous methods to find all records satisfying  $X$ , and check each for  $Y$

### ...Implementation of SELECT search

$\sigma$  (ESSN=123456789 AND PNO=10) (WORKS\_ON)

- **Conjunctive Selection using Composite Index**
- e.g. if an index exists on the composite key (ESSN, PNO), use the index directly
- Works if two or more attributes are used in an equality condition, and a composite index exists on those attributes

### ...Implementation of SELECT search

$\sigma$  (ESSN=123456789 AND PNO=10) (WORKS\_ON)

- **Conjunctive Selection by Pointer Intersection**
- Works if selection condition is a conjunction of equality conditions, and an index, with record pointers, exists for each attribute in the condition
- Strategy: index on each attribute value to find pointers to those records and take intersection of pointers
- Other conditions can also be tested within the intersection result.
- (Pointers must be to records, not blocks)

## Implementations for SELECT

$\sigma_{(ssn=123456789)}(EMPLOYEE)$

$\sigma_{(DNUMBER > 5)}(DEPARTMENT)$

$\sigma_{(DNO = 5)}(EMPLOYEE)$

- For these we can check whether an index, or hash, exists for the selection attribute, otherwise use brute force
- Optimization is needed most when a conjunction is involved

## Using Selectivity

- To optimize a query it might be helpful to know how many records would satisfy a condition such as  $X = \text{'Ottawa'}$
- If  $X$  is a key a selectivity estimate would be  $1 / |r(R)|$
- For a non-key attribute value  $X$  with  $i$  distinct values a selectivity estimate is  $|r(R)| / i / |r(R)| = 1/i$
- So database might keep track of the number of distinct values of an attribute (statistics about the data)

## Disjunctive Conditions (OR)

$\sigma_{(DNO = 5 \text{ OR } SALARY > 30000 \text{ OR } SEX = F)}(EMPLOYEE)$

- Harder to deal with
- Little optimization can be done because the condition requires union of records
- If any one attribute does not have an index, we are forced to use brute force approach
- Only when every OR condition attribute has an index can we optimize

## Implementing the JOIN operation

$EMPLOYEE \bowtie_{(DNO = DNUMBER)} DEPARTMENT$

$DEPARTMENT \bowtie_{(MGRSSN = SSN)} EMPLOYEE$

- JOIN is one of the most time-consuming operations in query processing
- Most joins are equi-joins (natural joins) so we only examine these

$R1 \bowtie_{(X=Y)} R2$

### ...Implementing the JOIN operation

EMPLOYEE  $\otimes$  (DNO = DNUMBER) DEPARTMENT

DEPARTMENT  $\otimes$  (MGRSSN = SSN) EMPLOYEE

- Brute Force (Nested Loop)
- For each t in EMPLOYEE {  
    For each s in DEPARTMENT {  
        if (t[DNO] = s[DNUMBER])  
            result := result union (t union s)  
    }  
}

### ...Implementing the JOIN operation

EMPLOYEE  $\otimes$  (DNO = DNUMBER) DEPARTMENT

- Using index on join attribute
- e.g. if index exists on DNO in EMPLOYEE
- Retrieve each record t from DEPARTMENT and use index to test for EMPLOYEE records which match t[DNUMBER]

### ...Implementing the JOIN operation

EMPLOYEE  $\otimes$  (DNO = DNUMBER) DEPARTMENT

- **Using Sort-merge JOIN**
- **If EMPLOYEE is physically sorted by DNO and DEPARTMENT sorted by DNUMBER**
- **Each file is scanned and matched against the other**
- **This method is very efficient since each file is only scanned once**
- **$O(n)$  for  $n$  records in largest file**

### ...Implementing the JOIN operation

EMPLOYEE  $\otimes$  (DNO = DNUMBER) DEPARTMENT

- **Using Hash JOIN**
- **If EMPLOYEE and DEPARTMENT are both hashed to the same file using the same hash function**
- **Requires a single pass through each file to hash into the same buckets**
- **Records from each file with equal join attributes will hash to the same bucket**
- **$O(n)$  for  $n$  records in largest file**

### Cost of JOIN operation

- Although we collect records with equal join attributes, it's actually disk block reads that count
- Strategy matters -consider brute force nested loop with

**DEPARTMENT**

**rD = 50 records**

**bD = 10 disk blocks**

**EMPLOYEE**

**rE = 5000 records**

**bE = 2000 disk blocks**

**nb = 6 blocks buffer in main memory**

### Cost of JOIN operation

**EMPLOYEE  $\otimes$  (DNO = DNUMBER) DEPARTMENT**

- For each t in EMPLOYEE {  
    For each s in DEPARTMENT {  
        if (t[DNO] = s[DNUMBER])  
            result := result union (t union s)  
    }  
}
- Strategy: read nb - 1 blocks from outer loop, then read through the entire inner loop file one block at a time
- Repeat for next nb-1 outer file blocks, etc.



### Cost of JOIN operation

EMPLOYEE  $\bowtie$  (DNO = DNUMBER) DEPARTMENT

- For each t in EMPLOYEE {  
    For each s in DEPARTMENT {  
        if (t[DNO] = s[DNUMBER])  
            result := result union (t union s)  
    }  
}
- Number blocks read for EMPLOYEE =  $b_E = 2000$
- Number of times nb-1 blocks are read  
    =  $b_E / (nb-1) = 2000 / 5 = 400$
- Total number of blocks read for inner file  
    =  $b_D * b_E / (nb-1) = 10 * 400 = 4000$
- Total number block access  
    =  $b_E + b_E / (nb-1) * b_D = 2000 + 2000 / 5 * 10 = 6000$

### Cost of JOIN operation

EMPLOYEE  $\bowtie$  (DNO = DNUMBER) DEPARTMENT

- For each s in DEPARTMENT {  
    For each t in EMPLOYEE {  
        if (t[DNO] = s[DNUMBER])  
            result := result union (t union s)  
    }  
}
- Number blocks read for DEPARTMENT =  $b_D = 10$
- Number of times nb-1 blocks are read  
    =  $b_D / (nb-1) = 10 / 5 = 2$
- Total number of blocks read for inner file  
    =  $b_E * b_D / (nb-1) = 2000 * 2 = 4000$
- Total number block access  
    =  $b_D + b_D / (nb-1) * b_E = 10 + 10 / 5 * 2000 = 4010$

- **Implementation Rule:** Use the smaller file in the outer loop of a nested-loop JOIN

### ...Implementing the JOIN operation

DEPARTMENT  $\otimes$  (MGRSSN = SSN) EMPLOYEE

- **Assume secondary indices exist on SSN in EMPLOYEE and MGRSSN in DEPARTMENT**
- **index level xSSN = 4, index level xMGRSSN = 2**
- **OPTION 1 Retrieve each record e from EMPLOYEE and use index to test DEPARTMENT for records whose MGRSSN match e[SSN]**
- **OPTION 2 Retrieve each record d from DEPARTMENT and use index to test EMPLOYEE for records whose SSN match d[MGRSSN]**

### ...Implementing the JOIN operation

DEPARTMENT  $\bowtie$  (MGRSSN = SSN) EMPLOYEE

- **OPTION 1** Retrieve each record  $e$  from EMPLOYEE and use index to test DEPARTMENT for records whose MGRSSN match  $e[SSN]$
- total block accesses is approx.  
 $= bE + (rE * (xMGRSSN + 1))$   
 $= 2000 + (5000 * (2 + 1)) = 17,000$  block accesses

**Note:** most EMPLOYEE records (4050 out of 5000) won't match join condition

### ...Implementing the JOIN operation

DEPARTMENT  $\bowtie$  (MGRSSN = SSN) EMPLOYEE

- **OPTION 2** Retrieve each record  $d$  from DEPARTMENT and use index to test EMPLOYEE for records whose SSN match  $d[MGRSSN]$
- total block accesses is approx.  
 $= bD + (rD * (xSSN + 1))$   
 $= 10 + (50 * (4 + 1)) = 260$  block accesses

**Note:** every DEPARTMENT record will match some employee record

- **Implementation Rule:** For join using index structure
- Use either the smaller file, or the file that has a match for every record (high selection factor) in the outer loop
- If can be worthwhile for the database to create an index just to do the join

### Optimization based on Heuristics

- **Heuristic:** Apply SELECT and PROJECT operations before applying JOIN
- Both SELECT and PROJECT reduce the size of the files the JOIN must work with, and hence the result

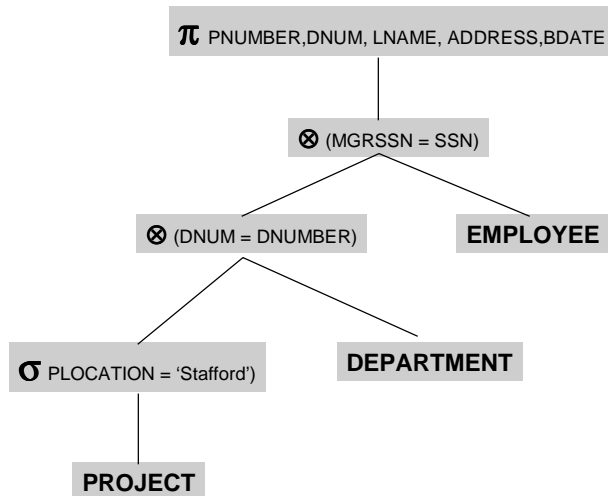
## Example Query

- For every project in 'Stafford' retrieve the project number, controlling department number and name, address and birthdate of the department manager

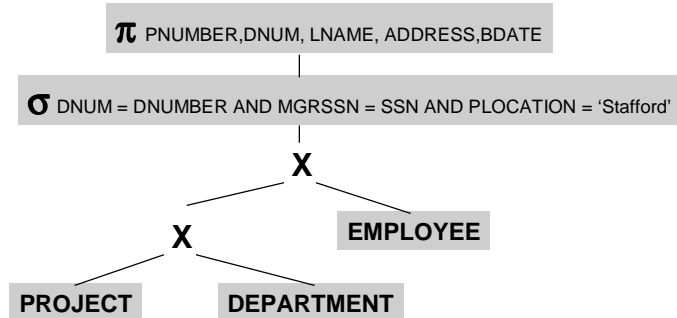
$\pi$  PNUMBER,DNUM, LNAME, ADDRESS,BDATE ( ( (   
 $\sigma$  PLOCATION = 'Stafford') (PROJECT))  $\otimes$  (DNUM = DNUMBER(   
 (DEPARTMENT)  $\otimes$  (MGRSSN = SSN( (EMPLOYEE))

**SELECT** PNUMBER, DNUM, LNAME, ADDRESS, BDATE  
**FROM** PROJECT, DEPARTMENT, EMPLOYEE  
**WHERE** DNUM=DNUMBER **AND**  
 MGRSSN=SSN **AND**  
 PLOCATION = 'Stafford'

## Query Tree (From relational Algebra)



## Equivalent Canonical Query Tree



- Canonical Tree is easy to generate but is very inefficient
- For PROJECT(100 tuples by 100 bytes)  
DEPARTMENT(50 tuples by 50 bytes)  
EMPLOYEE(5000 tuples by 150 bytes)
- Cartesian product would be: 10,000,000 by 300

© Louis D. Nel 1996

95.305

Query Processing 12 -

43

## Example Query

- Find the last names of employees born after 1957 who work on project "Aquarius"
- **SELECT** LNAME  
**FROM** EMPLOYEE, WORKS\_ON, PROJECT  
**WHERE** PNAME = 'Aquarius' **AND**  
PNUMBER = PNO **AND**  
ESSN = SSN **AND**  
BDATE > 'DEC-31-1957'

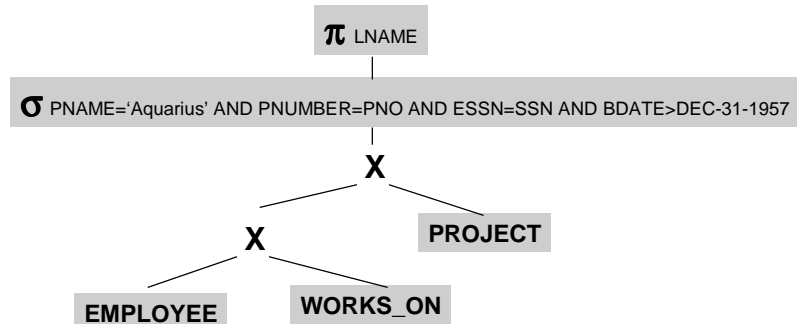
© Louis D. Nel 1996

95.305

Query Processing 12 -

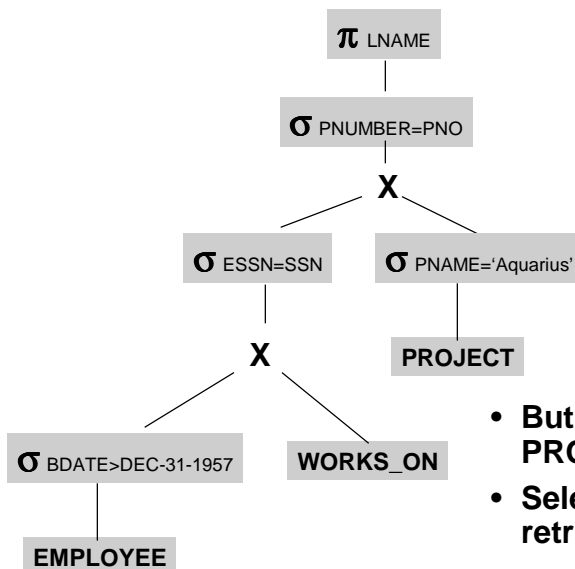
44

## Equivalent Canonical Query Tree



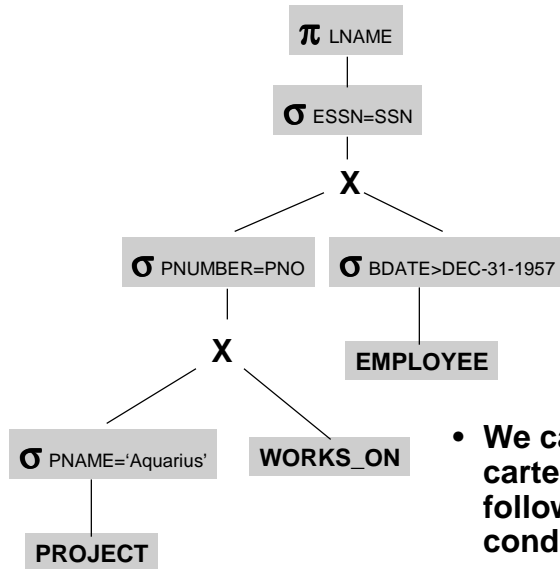
- But we need only one project record -for 'Aquarius'
- We need only those employees with BDATE>DEC-31-1957
- We can produce an equivalent, more efficient, tree

## Equivalent Query Tree



- But PNUMBER is key in PROJECT
- Select on PROJECT will retrieve a single record

## Equivalent Query Tree



- We can replace and cartesian product followed by a join condition with a JOIN

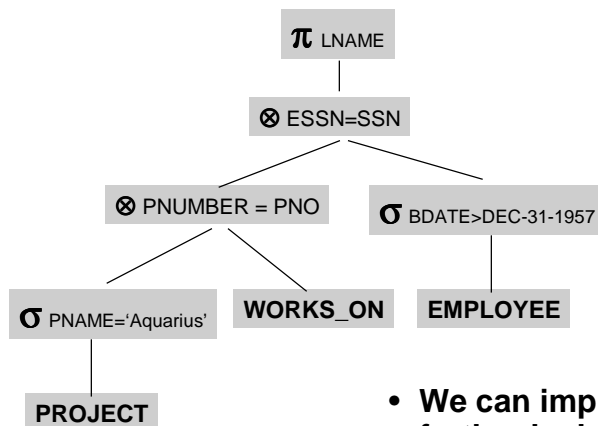
© Louis D. Nel 1996

95.305

Query Processing 12 -

47

## Equivalent Query Tree



- We can improve things further by keeping only the attributes we need
- i.e. do projects early

© Louis D. Nel 1996

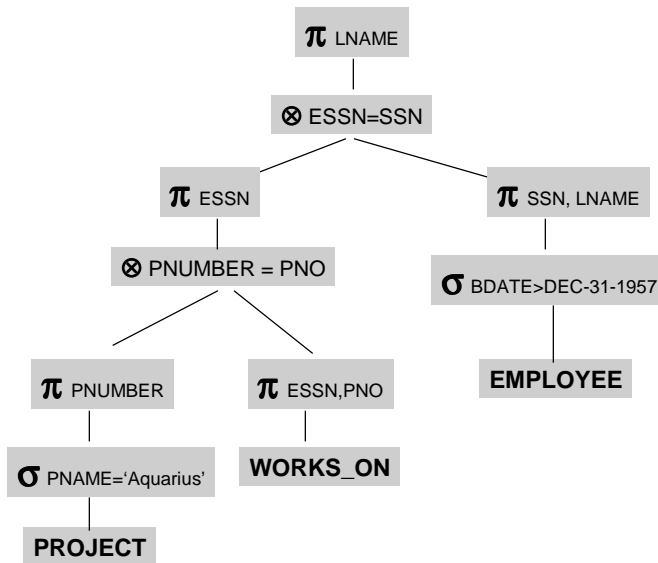
95.305

Query Processing 12 -

48



## Equivalent Final Query Tree



© Louis D. Nel 1996

95.305

Query Processing 12 -

49

## Optimization based on Cost Estimates

- Optimizer can also compare various execution strategies by computing an estimated cost and picking the best
- Optimizer should not spend more time estimating costs than it would take to execute the actual query

© Louis D. Nel 1996

95.305

Query Processing 12 -

50

## Query Execution Costs

- **Access Cost**: cost of searching for, reading, and writing disk blocks (Main cost)
- **Storage Cost**: cost of storing any temporary files generated by the execution strategy
- **Computation Cost**: cost of performing in-memory operations on data
- **Communication Cost**: Cost of shipping the query and its results from the database to client

## Cost meta-data

- To be able to estimate cost of queries we need to keep statistics about the actual data (meta-data)
- meta-data includes
  - r -> number of tuples in table
  - b -> number of blocks in table
  - bfr -> blocking factor (records per block)
  - x -> number of levels of index
  - bl1 -> number of first level index blocks
  - d -> number of distinct data values for index attr.
  - $s = r/d$  -> average number of index tuples with same value (selection cardinality)

### Cost Estimates for SELECT Strategies

- Linear Search: Cost  $\sim b/2$
- Binary Search: Cost  $\sim \log_2(b) + s/bfr - 1$
- Using Primary Key or Hash Key  
Cost  $\sim x + 1$ , (Cost  $\sim 1$  for Hashing)
- Using Ordering Index for multiple records  
Cost  $\sim x + b/2$
- Using a clustering index Cost  $\sim x + s/bfr$
- Using a Secondary B+ tree index  
Cost  $\sim x + s$  (for equality test)  
Cost  $\sim x + bl/2 + r/2$  (for range query)
- Conjunctive Selection  
Same is Linear, Binary or B+ tree
- Conjunctive with composite index  
Same as Primary Index or B+ tree