

Module 13

Transaction Processing Concepts

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 1

95.305

Objectives

- **Learn about issues transaction processing, concurrency and crash recovery**

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 2

95.305

Topics

- **Read and Write operations**
- **Concurrency**
- **Transactions**
- **Schedules and Recovery**

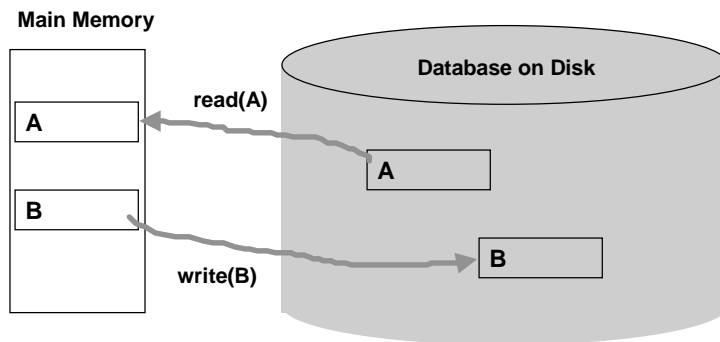
References

- **Elmasri & Navathe Chapters 17**

Observations

- **Computers, sometimes fail**
- **Databases are designed to survive computer crashes**
- **Databases must move from consistent state to consistent state**
- **Databases are typically multi-user and it's desirable to interleave transactions from different users**

Storage Model



read(A)
if block(A) not in main
memory
read A from disk to main
memory
 $x := A$

write(B)
if block(B) not in main
memory
read B from disk to main
memory
B in Block(B) $:= x$

Observations

- **Both read() and write() primitives may require disk block reads**
- **Neither operation specifically requires a disk block write**
- **Modified data block will eventually be written out to disk -but perhaps long after the transaction**
- **But, if system crashes between write() primitive and writing block back to disk, the modification is lost**

- **From the database point of view, transactions are user submissions which consist of**

read(X)

...

computation stuff

...

write(X)

- **Transactions submitted by various users may execute concurrently and may access the same data**
- **Different, concurrent transactions can interact in a bad way**

Simple Transactions

**T1: transfer \$100 from
savings to chequing
account**

```
read(X)
X := X-N
write(X)
read(Y)
Y := Y+N
write(Y)
```

**T2: deposit \$50 into
savings account**

```
read(X)
X:=X+M
write(X)
```

Serial Execution

		SavAcct	CheqAcct	
		\$1000	\$2000	{ \$1000 + \$2000 = \$3000 }
T1: transfer \$50 from savings to chequing account	read(X) X := X-N write(X) read(Y) Y := Y+N write(Y)	\$950		
			\$2050	{ \$950 + \$2050 = \$3000 }
T2: deposit \$150 into savings account	read(X) X:=X+M write(X)	\$1100		
		\$1100	\$2050	

Concurrent Execution: Lost update problem			
T1: transfer \$50 from savings to cheq. acct	T2: deposit \$150 into savings acct	SavAcct \$1000	CheqAcct \$2000
read(X)			
X := X-N			
	read(X)		
	X:=X+M		
write(X)		\$950	
read(Y)			
	write(X)	\$1150	
Y := Y+N			\$2050
write(Y)			
		\$1150	\$2050
© Louis D. Nel 1996	95.305	Transaction Processing 13 - 11	

Concurrent Execution: Dirty Read Problem			
T1: transfer \$50 from savings to cheq. acct	T2: deposit \$150 into savings acct	SavAcct \$1000	CheqAcct \$2000
read(X)			
X := X-N			
write(X)		\$950	
	read(X)		
	X:=X+M		
	write(X)	\$1100	
read(Y)			
t1 crash		\$1000	
		\$1000	\$2000
		T1: crashes and the database restores values written by T1 to their starting values	
© Louis D. Nel 1996	95.305	Transaction Processing 13 - 12	

Concurrent Execution: Incorrect Summary Problem

T1: transfer \$50 from savings to cheq. acct	T3: add balances of sav. & cheq. acct	SavAcct \$1000	CheqAcct \$2000	
read(X)				
X := X-N				
write(X)		\$950		
	read(X)			
	sum := sum + X			
	read(Y)			
	sum := sum + Y			Sum
	write(sum)			\$2950
read(Y)				
Y := Y+N				
write(Y)		\$950	\$2050	(\$3000)
© Louis D. Nel 1996	95.305	Transaction Processing 13 - 13		

- When a transaction is submitted the DBMS must make sure that either
 - a) All operations of the transaction are completed successfully
 - b) The transaction has no effect whatsoever
- Interleaved transactions may not appear to fail but may lead to inconsistencies

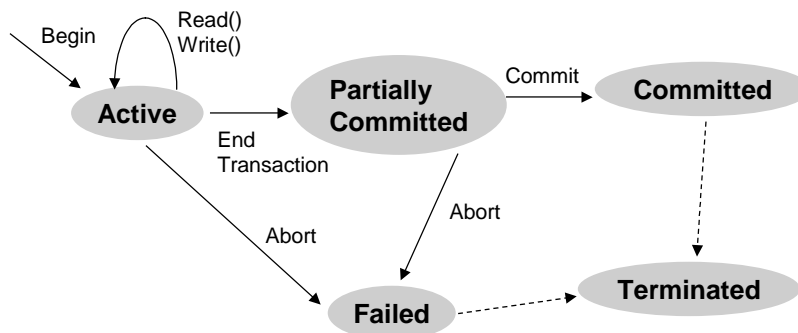
Common Failures

- 1) Computer Crash
- 2) Transaction Crash -operation in a transaction crashes (e.g. dividing account balance by zero)
- 3) Local error: data not found for transaction to proceed
- 4) Concurrency control enforcement: a transaction is terminated because it conflicts with a currently executing one

The database must maintain sufficient information to recover from these errors

Transaction States

- Recovery Mechanism keeps track of a transaction's state



- What happens if a crash occurs in the Partially Committed state

Transaction Journal

- To be able to recover from failures the DBMS maintains a log, or journal
- Log is kept in database and written to disk before transaction is deemed committed
- Log entries:
 - start(TransID)
 - write(TransID Dataltem oldValue, newValue)
 - read(TransID, Dataltem)
 - commit(TransID)
 - abort(TransID)
- Assumptions: transactions don't nest, all permanent changes to database are done through transactions

- Transaction reaches its committed state when
 - all transaction operations have completed successfully

AND

 - commit(TransID) has been written to the log (on disk)

Serial Execution			
	SavAcct	CheqAcct	Log
	\$1000	\$2000	start(T1)
T1: transfer \$50 from savings to chequing account	read(X)		read(T1, X)
	X := X-N		write(T1, X, 1000, 950)
	write(X)		read(T1, Y)
	\$950		write(T1, Y, 2000, 2050)
	read(Y)		commit(T1)
	Y := Y+N		
	write(Y)	\$2050	start(T2)
T2: deposit \$150 into savings account	read(X)		read(T1, X)
	X:=X+M		write(T1, X, 950, 1100)
	write(X)		commit(T2)
	\$1100		
	\$1100	\$2050	
© Louis D. Nel 1996	95.305	Transaction Processing 13 -	19

Crash Recovery	
<ul style="list-style-type: none"> • can use a redo(TranID) primitive • redo(TransID) “set the value of all data items updated by transaction TransID to their new Value” • redo() primitive must be idempotent: executing it several times must be equivalent to executing it once • Transaction TransID needs to be redone if a crash occurs and both start(TransID) and commit(TransID) appears in the log • Log could be periodically checkpoint to ensure rollback need not go to far back 	
© Louis D. Nel 1996	95.305 Transaction Processing 13 - 20

Desirable Transaction Properties

- **Atomic**: transactions are either performed entirely or not at all
- **Consistency Preservation**: correct execution must take database from consistent state to consistent state
- **Isolation**: a transaction should not make its updates visible to other transactions until it has been committed (precludes concurrency)
- **Durability**: once a transaction has been committed its effects must not be lost due to subsequent failure

Buffer Management

- Log entries are written to disk before operations re executed
- Wasteful to do block write for every log entry
- Protocol: before commit(T) is written, all log records pertaining to T must be written to disk
- Before a data block is written all log records must have been written to disk
- Notice this puts a strain on virtual memory -the host OS might not want to swap pages this way

Concurrency

- **Concurrency:** allowing several transactions to execute in an interleaved way

Interleaving OK : does not produce inconsistency

T1: transfer \$50 from savings to cheq. acct	T2: deposit \$150 into savings acct	SavAcct \$1000	CheqAcct \$2000
read(X)			
X := X-N			
write(X)		\$950	
	read(X)		
	X:=X+M		
	write(X)	\$1100	
read(Y)			
Y := Y+N			
write(Y)			\$2050
		\$1100	\$2050

Bad Interleaving: produces inconsistencies

T1: transfer \$50 from savings to cheq. acct	T2: deposit \$150 into savings acct	SavAcct \$1000	CheqAcct \$2000
read(X)			
X := X-N			
	read(X)		
	X:=X+M		
write(X)		\$950	
read(Y)			
	write(X)	\$1150	
Y := Y+N			\$2050
write(Y)			
		\$1150	\$2050

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 25

Serializable Schedules

- An interleaving of transactions is OK if it is equivalent to some serial schedule

T1: transfer \$50 from savings to cheq. acct	T2: deposit \$150 into savings acct		T1: transfer \$50 from savings to cheq. acct	T2: deposit \$150 into savings acct
read(X)			read(X)	
X := X-N			X := X-N	
write(X)			write(X)	
	read(X)	=	read(Y)	
	X:=X+M		Y := Y+N	
	write(X)		write(Y)	
read(Y)				read(X)
Y := Y+N				X:=X+M
write(Y)				write(X)

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 26

Serializable Schedules

- An interleaving of transactions is OK if it is equivalent to some serial schedule

T1: transfer \$50 from
savings to cheq.
acct

read(X)

X := X-N

write(X)

read(Y)

Y := Y+N

write(Y)

T2: deposit
\$150 into
savings acct

read(X)

X:=X+M

write(X)

≠

T1: transfer \$50 from
savings to cheq.
acct

read(X)

X := X-N

write(X)

read(Y)

Y := Y+N

write(Y)

T2: deposit
\$150 into
savings acct

read(X)

X:=X+M

write(X)

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 27

- Idea: DBMS can allow interleaving as long as it remains “convinced” that a serial equivalent schedule exists
- i.e. that the interleaving is equivalent to some serial schedule

© Louis D. Nel 1996

95.305

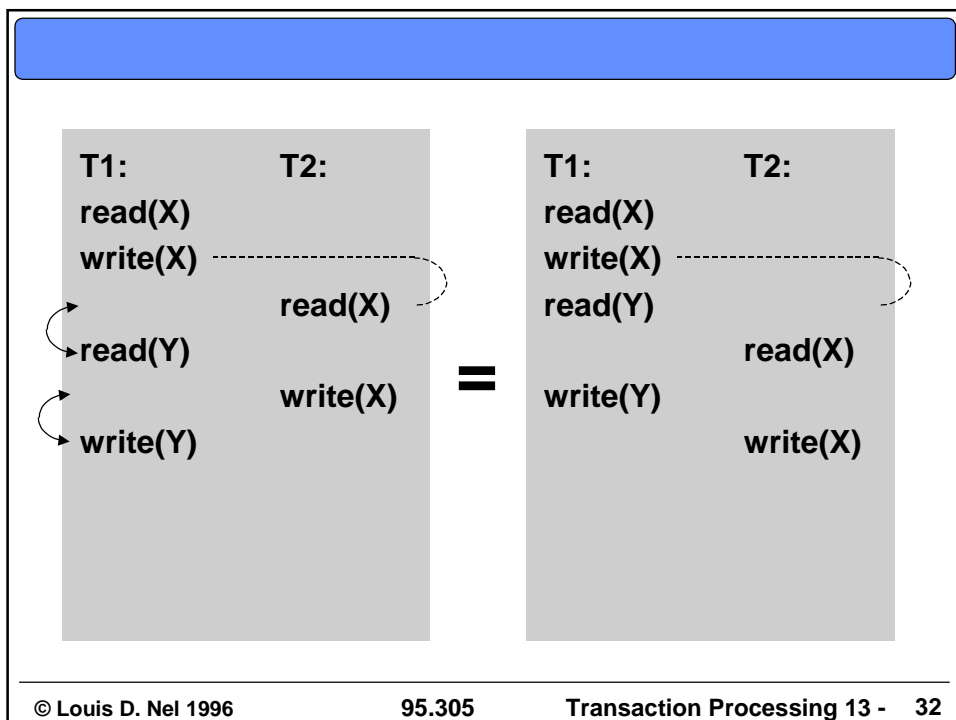
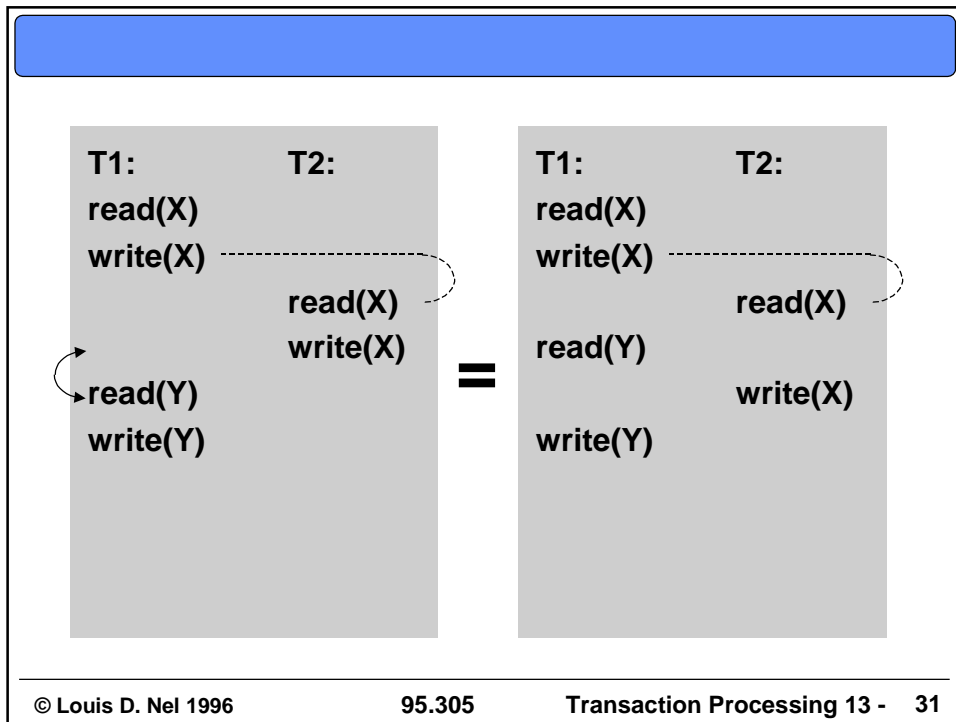
Transaction Processing 13 - 28

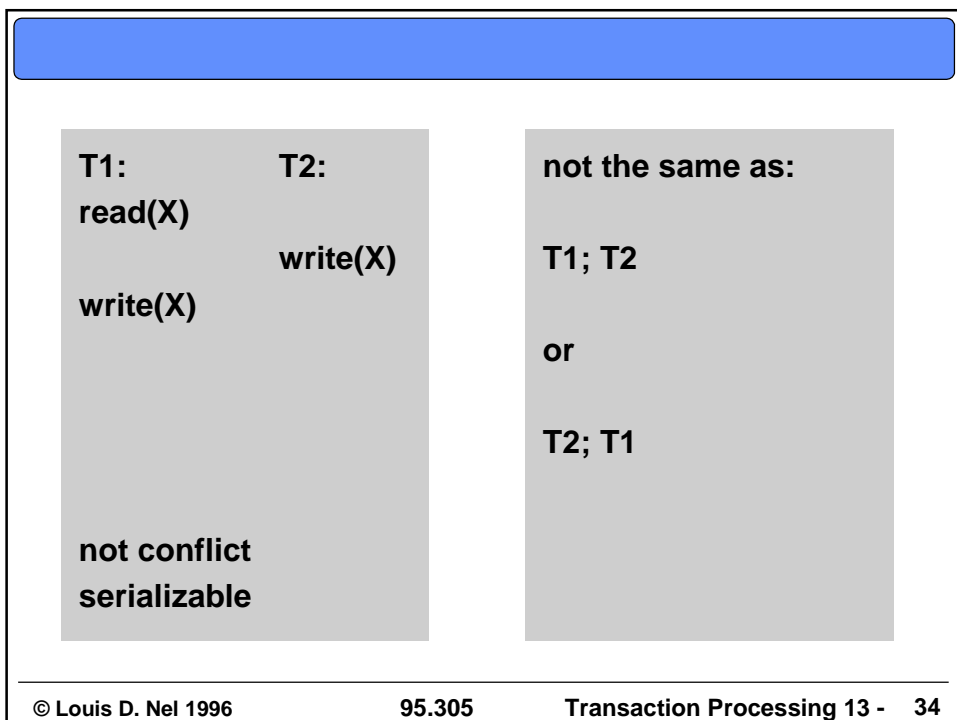
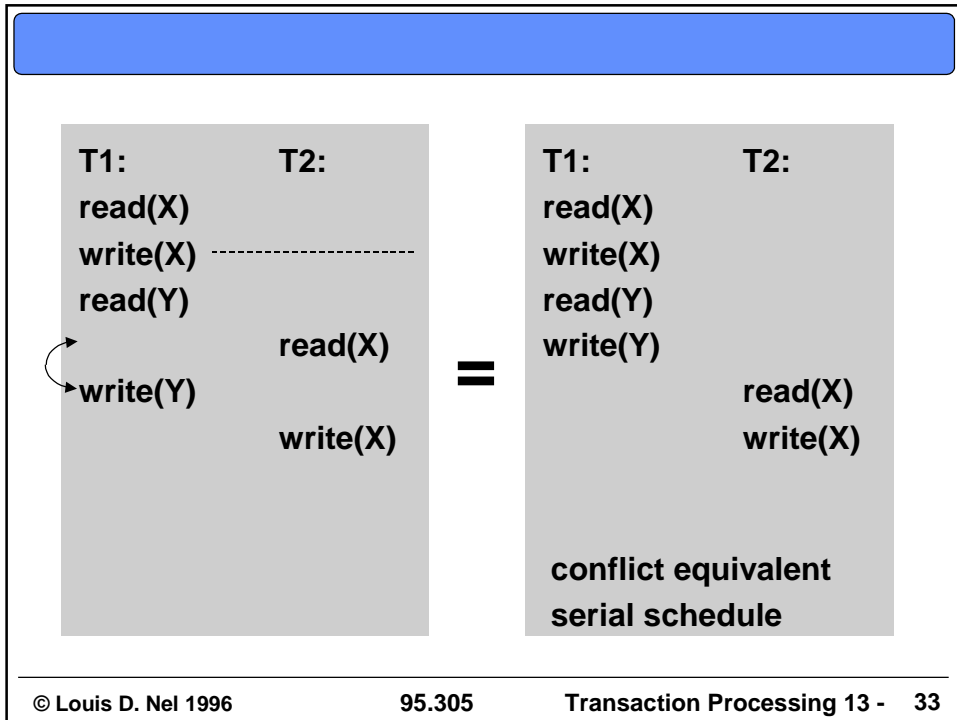
Simple Transaction Model

<p>read(X)</p> <p>...</p> <p>stuff</p> <p>...</p> <p>write(X)</p>	<p>=</p>	<p>read(X)</p> <p>write(X)</p>
---	----------	--------------------------------

Instruction Conflicts

<p>T1:</p> <p>read(X)</p> <p>read(X)</p> <p>write(X)</p> <p>write(X)</p>	<p>T2:</p> <p>read(X)</p> <p>write(X)</p> <p>read(X)</p> <p>write(X)</p>	<p>order doesn't matter</p> <p>order matters</p> <p>order matters</p> <p>order matters</p>
<p>read(X)</p> <p>read(X)</p> <p>write(X)</p> <p>write(X)</p>	<p>read(Y)</p> <p>write(Y)</p> <p>read(Y)</p> <p>write(Y)</p>	<p>order doesn't matter</p> <p>"</p> <p>"</p> <p>"</p>





Not serializable by equivalent

T1:
 read(X)
 i := X-50
 write(X)

T2:

 read(Y)
 j := Y-10
 write(Y)

read(Y)
 k := Y+50
 write(Y)

read(X)
 l := X+10
 write(X)

T1:
 read(X)
 i := X-50
 write(X)
 read(Y)
 k := Y+50
 write(Y)

T2:

 read(Y)
 j := Y-10
 write(Y)
 read(X)
 l := X+10
 write(X)

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 35

- Serializability is a sufficient, but not necessary, condition to prevent interleaving inconsistencies

© Louis D. Nel 1996

95.305

Transaction Processing 13 - 36

- **Serializability is too difficult to test for in practice**
- **Alternative is to determine what actions a transaction can take which will ensure that serializability results**
- **Transactions follow protocols which allows them to interleave with assurance of consistency**

- **Example protocol, or alternative, is to lock the data values with a mutual exclusion semaphore**
- **This ensures that only one transaction has access to data (prevent read and write by others)**
- **Problem: this can lead to deadlock and starve transactions which mutually need each others data**