

## **Chapter 18**

---

### **Query Processing**

### **Transparencies**

## **Chapter - Objectives**

---

- u **Objectives of query processing and optimization.**
- u **Static versus dynamic query optimization.**
- u **How a query is decomposed and semantically analyzed.**
- u **How to create a R.A.T. to represent a query.**
- u **Rules of equivalence for RA operations.**
- u **How to apply heuristic transformation rules to improve efficiency of a query.**

## **Chapter - Objectives**

---

- u **Types of database statistics required to estimate cost of operations.**
- u **Different strategies for implementing selection.**
- u **How to evaluate cost and size of selection.**
- u **Different strategies for implementing join.**
- u **How to evaluate cost and size of join.**
- u **Different strategies for implementing projection.**
- u **How to evaluate cost and size of projection.**

3

## **Introduction**

---

- u **In network and hierarchical DBMSs, low-level procedural query language is generally embedded in high-level programming language.**
- u **Programmer's responsibility to select most appropriate execution strategy.**
- u **With declarative languages such as SQL, user specifies what data is required rather than how it is to be retrieved.**
- u **Relieves user of knowing what constitutes good execution strategy.**

5

## **Introduction**

---

- u **Also gives DBMS more control over system performance.**
- u **Two main techniques for query optimization:**
  - **heuristic rules that order operations in a query.**
  - **comparing different strategies based on relative costs, and selecting one that minimizes resource usage.**
- u **Disk access tends to be dominant cost in query processing for centralized DBMS.**

6

## **Query Processing**

---

**Activities involved in retrieving data from the database.**

- u **Aims of QP:**
  - **transform query written in high-level language (e.g. SQL), into correct and efficient execution strategy expressed in low-level language (implementing RA);**
  - **execute the strategy to retrieve required data.**

7

## **Query Optimization**

---

**Activity of choosing an efficient execution strategy for processing query.**

- u As there are many equivalent transformations of same high-level query, aim of QO is to choose one that minimizes resource usage.**
- u Generally, reduce total execution time of query.**
- u May also reduce response time of query.**
- u Problem computationally intractable with large number of relations, so strategy adopted is reduced to finding near optimum solution.**

8

## **Example 18.1 - Different Strategies**

---

**Find all Managers that work at a London branch.**

```
SELECT *  
FROM staff s, branch b  
WHERE s.bno = b.bno AND  
(s.position = 'Manager' AND b.city = 'London');
```

9

### Example 18.1 - Different Strategies

---

u Three equivalent RA queries are:

(1)  $S_{(\text{position}='Manager')} \cup (city='London') \cup (staff.bno=branch.bno)$   
(Staff X Branch)

(2)  $S_{(\text{position}='Manager')} \cup (city='London') ($   
Staff  $\bowtie_{staff.bno=branch.bno}$  Branch)

(3)  $(S_{\text{position}='Manager'}(Staff)) \bowtie_{staff.bno=branch.bno}$   
 $(S_{city='London'}(Branch))$

10

### Example 18.1 - Different Strategies

---

u Assume:

- 1000 tuples in Staff; 50 tuples in Branch;
- 50 Managers; 5 London branches;
- No indexes or sort keys;
- Results of any intermediate operations stored on disk;
- Cost of the final write is ignored;
- Tuples are accessed one at a time.

11

### **Example 18.1 - Cost Comparison**

---

u **Cost (in disk accesses) are:**

**(1)  $(1000 + 50) + 2*(1000 * 50) = 101\ 050$**

**(2)  $2*1000 + (1000 + 50) = 3\ 050$**

**(3)  $1000 + 2*50 + 5 + (50 + 5) = 1\ 160$**

u **Cartesian product and join operations are much more expensive than selection, and third option significantly reduces size of relations being joined together.**

12

### **Phases of Query Processing**


---

u **QP has four main phases:**

- decomposition (consisting of parsing and validation)**
- optimization**
- code generation**
- execution.**

13

## Phases of Query Processing



14

## Dynamic versus Static Optimization

- u Advantages of static QO are removal of runtime overhead, and more time to find optimum strategy.
- u Disadvantages arise from fact that chosen execution strategy may no longer be optimal when query is run.
- u Could use a hybrid approach to overcome this.

16

## **Query Decomposition**

---

- u **Aims are to transform high-level query into RA query and check that query is syntactically and semantically correct.**
- u **Analyze query lexically and syntactically using compiler techniques.**
- u **Verify relations and attributes exist.**
- u **Verify operations are appropriate for object type.**

17

## **Analysis - Example**

---

```
SELECT staff_no  
FROM staff  
WHERE position > 10;
```

- u **This query would be rejected on two grounds:**
  - **Staff\_No is not defined for Staff relation (should be Sno).**
  - **Comparison '>10' is incompatible with type Position, which is variable character string.**

19



## Analysis

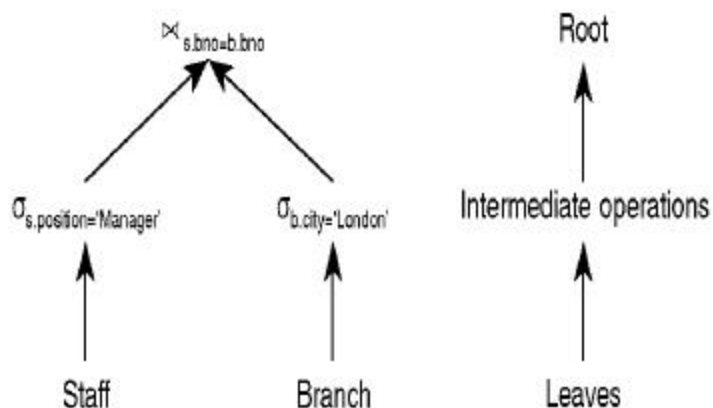
---

- u Finally, query transformed into some internal representation more suitable for processing.
- u Some kind of query tree is typically chosen, constructed as follows:
  - Leaf node created for each base relation.
  - Non-leaf node created for each intermediate relation produced by RA operation.
  - Root of tree represents query result.
  - Sequence is directed from leaves to root.

20

## Example 18.1 - R.A.T.

---



21

## **Semantic Analysis**

---

- u **Rejects normalized queries that are incorrectly formulated or contradictory.**
- u **Query is incorrectly formulated if components do not contribute to generation of result.**
- u **Query is contradictory if its predicate cannot be satisfied by any tuple.**
- u **Algorithms to determine correctness exist only for queries that do not contain disjunction and negation.**

23

## **Simplification**

---

- **Detects redundant qualifications,**
- **Eliminates common sub-expressions,**
- **Transforms query to semantically equivalent but more easily and efficiently computed form.**

29

## **Transformation Rules for RA Operations**

---

**Conjunctive selection operations can cascade into individual selection operations (and vice versa).**

$$S_{p \wedge q \wedge r}(R) = S_p(S_q(S_r(R)))$$

u **Sometimes referred to as cascade of selection.**

$$\begin{aligned} S_{\text{bno}='B3' \wedge \text{salary}>15000}(\text{Staff}) &= \\ S_{\text{bno}='B3'}(S_{\text{salary}>15000}(\text{Staff})) \end{aligned}$$

30

## **Transformation Rules for RA Operations**

---

**Commutativity of selection.**

$$S_p(S_q(R)) = S_q(S_p(R))$$

u **For example:**

$$\begin{aligned} S_{\text{bno}='B3'}(S_{\text{salary}>15000}(\text{Staff})) &= \\ S_{\text{salary}>15000}(S_{\text{bno}='B3'}(\text{Staff})) \end{aligned}$$

31

## **Transformation Rules for RA Operations**

---

**In a sequence of projection operations, only the last in the sequence is required.**

$$P_L P_M \dots P_N(R) = P_L(R)$$

u **For example:**

$$P_{lname} P_{bno}(Staff) = P_{lname}(Staff)$$

32

## **Transformation Rules for RA Operations**

---

**Commutativity of selection and projection.**

u **If predicate p involves only attributes in projection list, selection and projection operations commute:**

$$P_{A_1, \dots, A_m}(S_p(R)) = S_p(P_{A_1, \dots, A_m}(R))$$

where  $pl \supseteq \{A_1, A_2, \dots, A_m\}$

u **For example:**

$$P_{fname, lname}(S_{lname='Beech'}(Staff)) =$$
$$S_{lname='Beech'}(P_{fname, lname}(Staff))$$

33

## **Transformation Rules for RA Operations**

**Commutativity of theta-join (and Cartesian product).**

$$R \bowtie_p S = S \bowtie_p R$$

$$R \times S = S \times R$$

- u **Rule also applies to equi-join and natural join. For example:**

$$\text{Staff} \bowtie_{\text{staff.bno}=\text{branch.bno}} \text{Branch} =$$

$$\text{Branch} \bowtie_{\text{staff.bno}=\text{branch.bno}} \text{Staff}$$

34

## **Transformation Rules for RA Operations**

**Commutativity of selection and theta-join (or Cartesian product).**

- u **If selection predicate involves only attributes of one of join relations, selection and join (or Cartesian product) operations commute:**

$$s_p(R \bowtie_r S) = (s_p(R)) \bowtie_r S$$

$$s_p(R \times S) = (s_p(R)) \times S$$

$$\text{where } p \hat{=} \{A_1, A_2, \dots, A_n\}$$

35

## **Transformation Rules for RA Operations**

- u If selection predicate is conjunctive predicate having form  $(p \dot{\cup} q)$ , where  $p$  only involves attributes of  $R$ , and  $q$  only attributes of  $S$ , selection and theta-join operations commute as:

$$s_{p \dot{\cup} q}(R \bowtie_r S) = (s_p(R)) \bowtie_r (s_q(S))$$

$$s_{p \dot{\cup} q}(R \bowtie S) = (s_p(R)) \bowtie (s_q(S))$$

36

## **Transformation Rules for RA Operations**

- u For example:

$$s_{\text{position}='Manager' \dot{\cup} \text{city}='London'}(\text{Staff} \bowtie_{\text{staff.bno}=\text{branch.bno}} \text{Branch}) =$$

$$(s_{\text{position}='Manager'}(\text{Staff})) \bowtie_{\text{staff.bno}=\text{branch.bno}} (s_{\text{city}='London'}(\text{Branch}))$$

37

## **Transformation Rules for RA Operations**

**Commutativity of projection and theta-join (or Cartesian product).**

- u If projection list is of form  $L = L_1 \hat{\cup} L_2$ , where  $L_1$  only involves attributes of  $R$ , and  $L_2$  only involves attributes of  $S$ , provided join condition only contains attributes of  $L$ , projection and theta-join operations commute as:

$$P_{L_1 \hat{\cup} L_2}(R \bowtie_r S) = (P_{L_1}(R)) \bowtie_r (P_{L_2}(S))$$

38

## **Transformation Rules for RA Operations**

**Associativity of theta-join (and Cartesian product).**

- u Cartesian product and natural join are always associative:

$$(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$$

$$(R \times S) \times T = R \times (S \times T)$$

- u If join condition  $q$  involves attributes only from  $S$  and  $T$ , then theta-join is associative as follows:

$$(R \bowtie_p S) \bowtie_{q \cup r} T = R \bowtie_{p \cup r} (S \bowtie_q T)$$

44

## Transformation Rules for RA Operations

u For example:

$$\begin{aligned} & (\text{Staff} \bowtie_{\text{staff.sno}=\text{property\_for\_rent.sno}} \text{Property\_for\_Rent}) \\ & \bowtie_{\text{ono}=\text{owner.ono} \vee \text{staff.lname}=\text{owner.lname}} \text{Owner} = \\ & \text{Staff} \bowtie_{\text{staff.sno}=\text{property\_for\_rent.sno} \vee \text{staff.lname}=\text{owner.lname}} \\ & (\text{Property\_for\_Rent} \bowtie_{\text{ono}} \text{Owner}) \end{aligned}$$

45

## Example 18.3 Use of Transformation Rules


For prospective renters of flats, find properties that match requirements and owned by CO93.

```
SELECT p.pno, p.street
FROM renter r, viewing v, property_for_rent p
WHERE r.pref_type = 'Flat' AND
      r.rno = v.rno AND v.pno = p.pno AND
      r.max_rent >= p.rent AND
      r.pref_type = p.type AND p.ono = 'CO93';
```

46




### Example 18.3 Use of Transformation Rules



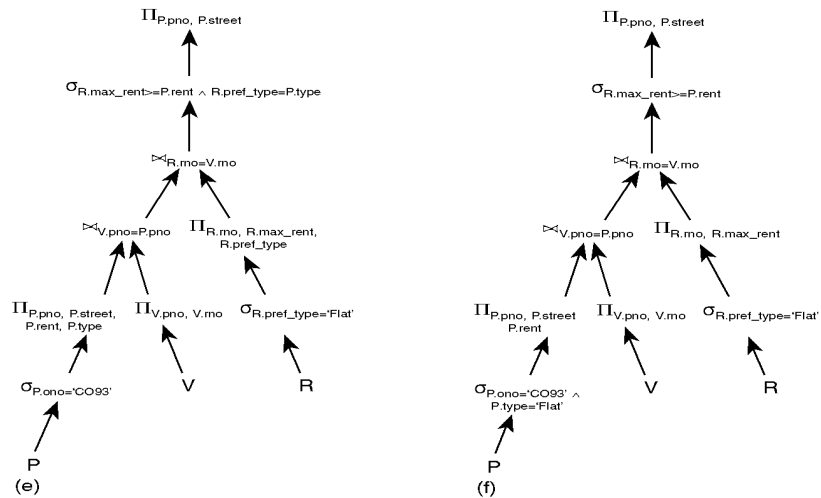
47

### Example 18.3 Use of Transformation Rules



48

### Example 18.3 Use of Transformation Rules



49

### Heuristical Processing Strategies

- u Perform selection operations as early as possible.
  - Keep predicates on same relation together.
- u Combine Cartesian product with subsequent selection whose predicate represents join condition into a join operation.
- u Use associativity of binary operations to rearrange leaf nodes so leaf nodes with most restrictive selection operations executed first.

50

## **Heuristical Processing Strategies**

---

- u **Perform projection as early as possible.**
  - **Keep projection attributes on same relation together.**
- u **Compute common expressions once.**
  - **If common expression appears more than once, and result not too large, store result and reuse it when required.**
  - **Useful when querying views, as same expression is used to construct view each time.**

51

## **Cost Estimation for RA Operations**

---

- u **Many different ways of implementing RA operations.**
- u **Aim of QO is to choose most efficient one.**
- u **Use formulae that estimate costs for a number of options, and select one with lowest cost.**
- u **Consider only cost of disk access, which is usually dominant cost in QP.**
- u **Many estimates are based on cardinality of the relation, so need to be able to estimate this.**

52

## **Database Statistics**

---

- u **Success of estimation depends on amount and currency of statistical information DBMS holds.**
- u **Keeping statistics current can be problematic.**
- u **If statistics updated every time tuple is changed, this would impact performance.**
- u **DBMS could update statistics on a periodic basis, for example nightly, or whenever the system is idle.**

53

## **Typical Statistics for Relation R**

---

**ntuples(R) - number of tuples in R.**

**bfactor(R) - blocking factor of R.**

**nblocks(R) - number of blocks required to store R:**

$$\mathbf{nblocks(R) = [ntuples(R)/bfactor(R)]}$$

54

### **Typical Statistics for Attribute A of Relation R**

**$\text{ndistinct}_A(R)$**  - number of distinct values that appear for attribute A in R.

**$\text{min}_A(R), \text{max}_A(R)$**

- minimum and maximum possible values for attribute A in R.

**$\text{SC}_A(R)$**  - selection cardinality of attribute A in R.

Average number of tuples that satisfy an equality condition on attribute A.

55

### **Statistics for Multilevel Index I on Attribute A**

**$\text{nlevels}_A(I)$**  - number of levels in I.

**$\text{nlfblocks}_A(I)$**  - number of leaf blocks in I.

56

## **Selection Operation**

---

- u **Predicate may be simple or composite.**
- u **Number of different implementations, depending on file structure, and whether attribute(s) involved are indexed/hashed.**
- u **Main strategies are:**
  - **Linear Search (Unordered file, no index).**
  - **Binary Search (Ordered file, no index).**
  - **Equality on hash key.**
  - **Equality condition on primary key.**

57

## **Equality of Hash Key**

---

- u **If attribute A is hash key, apply hashing algorithm to calculate target address for tuple.**
- u **If there is no overflow, expected cost is 1.**
- u **If there is overflow, additional accesses may be necessary.**

62

### **Equality Condition on Primary Key**

- u Can use primary index to retrieve single record satisfying condition.
- u Need to read one more block than number of index accesses, equivalent to number of levels in index, so estimated cost is:

$$\text{nlevels}_A(I) + 1$$

63

### **Inequality Condition on Primary Key**

- u Can first use index to locate record satisfying predicate ( $A = x$ ).
- u Provided index is sorted, records can be found by accessing all records before/after this one.
- u Assuming uniform distribution, would expect half the records to satisfy inequality, so estimated cost is:

$$\text{nlevels}_A(I) + [\text{nblocks}(R)/2]$$

64

## **Join Operation**

---

- u **Main strategies for implementing join:**
  - **Block Nested Loop Join.**
  - **Indexed Nested Loop Join.**
  - **Sort-Merge Join.**
  - **Hash Join.**

70

## **Estimating Cardinality of Join**

---

- u **Cardinality of Cartesian product is:**  
 **$\text{ntuples(R)} * \text{ntuples(S)}$**
- u **More difficult to estimate cardinality of any join as depends on distribution of values.**
- u **Worst case, cannot be any greater than this value.**

71



## **Estimating Cardinality of Join**

---

- u **If assume uniform distribution, can estimate for equi-joins with a predicate ( $R.A = S.B$ ) as follows:**
  - **If A is key of R:      $ntuples(T) \leq ntuples(S)$**
  - **If B is key of S:      $ntuples(T) \leq ntuples(R)$**
- u **Otherwise, could estimate cardinality of join as:**  
 **$ntuples(T) = SC_A(R) * ntuples(S)$  or**  
 **$ntuples(T) = SC_B(S) * ntuples(R)$**

72

## **Block Nested Loop Join**

---

- u **Simplest join algorithm is nested loop that joins two relations together a tuple at a time.**
- u **Outer loop iterates over each tuple in R, and inner loop iterates over each tuple in S.**
- u **As basic unit of reading/writing is a disk block, better to have two extra loops that process blocks.**
- u **Estimated cost of this approach is:**

$$nblocks(R) + (nblocks(R) * nblocks(S))$$

73

## **Indexed Nested Loop Join**

---

- u **If have index (or hash function) on join attributes of inner relation, can use index lookup.**
- u **For each tuple in R, use index to retrieve matching tuples of S.**
- u **Cost of scanning R is nblocks(R), as before.**
- u **Cost of retrieving matching tuples in S depends on type of index and number of matching tuples.**
- u **If join attribute A in S is PK, cost estimate is:**  
$$\text{nblocks(R)} + \text{ntuples(R)} * (\text{nlevels}_A(\text{I}) + 1)$$

75

## **Sort-Merge Join**

---

- u **For equi-joins, most efficient join is when both relations are sorted on join attributes.**
- u **Can look for qualifying tuples merging relations.**
- u **May need to sort relations first.**
- u **Now tuples with same join value are in order.**
- u **If assume join is M:N and each set of tuples with same join value can be held in database buffer at same time, then each block of each relation need only be read once.**

76

## **Sort-Merge Join**

---

- u **Cost estimate for the sort-merge join is:**

$$\text{nblocks(R)} + \text{nblocks(S)}$$

- u **If a relation has to be sorted, R say, add:**

$$\text{nblocks(R)} * [\log_2(\text{nblocks(R)})]$$

77

## **Hash Join**

---

- u **For natural or equi-join, hash join may be used.**
- u **Idea is to partition relations according to some hash function that provides uniformity and randomness.**
- u **Each equivalent partition should hold same value for join attributes, although it may hold more than one value.**
- u **Cost estimate of hash join as:**

$$3(\text{nblocks(R)} + \text{nblocks(S)}) + 2 * \text{max\_partitions}$$

78

## **Projection Operation**

---

- u **To implement projection need following steps:**
  - **Removal of attributes that are not required.**
  - **Elimination of any duplicate tuples produced from previous step. Only required if projection attributes do not include a key.**
- u **Two main approaches to eliminating duplicates:**
  - **Sorting**
  - **Hashing.**

79

## **Duplicate Elimination using Sorting**

---

- u **Sort tuples of reduced relation using all remaining attributes as sort key.**
- u **Duplicates will now be adjacent and can be removed easily.**
- u **Estimated cost of sorting is:**  
$$\text{nblocks(R)} * [\log_2(\text{nblocks(R)})].$$
- u **Combined cost is:**  
$$\text{nblocks(R)} + \text{nblocks(R)} * [\log_2(\text{nblocks(R)})]$$

81

## **Duplicate Elimination using Hashing**

---

- u **Two phases: partitioning and duplicate elimination.**
- u **In partitioning phase, for each tuple in R, remove unwanted attributes and apply hash function to combination of remaining attributes, and write reduced tuple to hashed value.**
- u **Two tuples that belong to different partitions are guaranteed not to be duplicates.**
- u **Estimated cost is:  $nblocks(R) + nb$**

82

## **Aggregate Operations**

---

**SELECT avg(salary)  
FROM staff;**

- u **To implement query, could scan entire Staff relation and maintain running count of number of tuples read and sum of all salaries.**
- u **Easy to compute average from these two running counts.**

85

## **Aggregate Operations**

---

```
SELECT avg(salary)  
FROM staff  
GROUP BY bno;
```

- u **For grouping queries, can use sorting or hashing algorithms similar to duplicate elimination.**
- u **Can estimate cardinality of result using estimates derived earlier for selection.**