Chapter 20

Distributed DBMSs - Advanced Concepts

Transparencies

Chapter - Objectives

- **Distributed transaction management.**
- **Distributed concurrency control.**
- **Distributed deadlock detection.**

Distributed Transaction Management

- **Distributed transaction accesses data stored at more than one location.**
- **Divided into a number of *sub-transactions,* one for each site that has to be accessed, represented by an *agent.***
- **Indivisibility of distributed transaction is still fundamental to transaction concept.**
- **DDBMS must also ensure indivisibility of each sub-transaction.**

Distributed Transaction Management

- **Thus, DDBMS must ensure:**
  - **synchronization of subtransactions with other local transactions executing concurrently at a site;**
  - **synchronization of subtransactions with global transactions running simultaneously at same or different sites.**
- **Global transaction manager (transaction coordinator) at each site, to coordinate global and local transactions initiated at that site.**

Coordination of Distributed Transaction

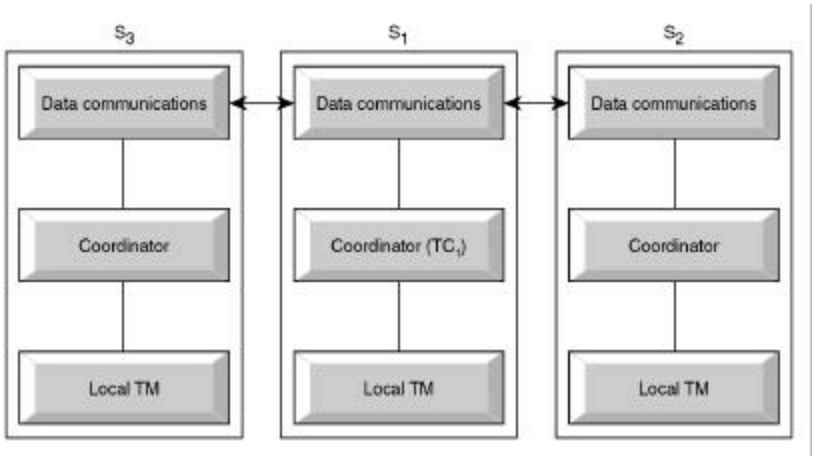

Distributed Locking

- **Look at four schemes:**
  - **Centralized locking**
  - **Primary Copy 2PL**
  - **Distributed 2PL**
  - **Majority Locking**

Centralized Locking

- **Single site that maintains all locking information.**
- **One lock manager for whole of DDBMS.**
- **Local transaction managers involved in global transaction request and release locks from lock manager.**
- **Or transaction coordinator can make all locking requests on behalf of local transaction managers.**
- **Advantage - easy to implement.**
- **Disadvantages - bottlenecks and lower reliability.**

Primary Copy 2PL

- **Lock managers distributed to a number of sites.**
- **Each lock manager responsible for managing locks for set of data items.**
- **For replicated data item, one copy is chosen as *primary copy,* others are *slave copies***
- **Only need to write-lock primary copy of data item that is to be updated.**
- **Once primary copy has been updated, change can be propagated to slaves.**

Primary Copy 2PL

- **Disadvantages - deadlock handling is more complex due; still a degree of centralization in system.**
- **Advantages - lower communication costs and better performance than centralized 2PL.**

Distributed 2PL

- **Lock managers distributed to every site.**
- **Each lock manager responsible for locks for data at that site.**
- **If data not replicated, equivalent to primary copy 2PL.**
- **Otherwise, implements a Read-One-Write-All (ROWA) replica control protocol.**

Distributed 2PL

- **Using ROWA protocol:**
  - **Any copy of replicated item can be used for read.**
  - **All copies must be write-locked before item can be updated.**
- **Disadvantages - deadlock handling more complex; communication costs higher than primary copy 2PL.**

Majority Locking

- **Extension of distributed 2PL.**
- **To read or write data item replicated at *n* sites, sends a lock request to more than half the *n* sites where item is stored.**
- **Transaction cannot proceed until majority of locks obtained.**
- **Overly strong in case of read locks.**

Distributed Deadlock

- More complicated if lock management is not centralized.
- Local Wait-for-Graph (LWFG) may not show existence of deadlock.
- May need to create GWFG, union of all LWFGs.
- Look at three schemes:
  - Centralized Deadlock Detection
  - Hierarchical Deadlock Detection
  - Distributed Deadlock Detection.

Example - Distributed Deadlock

- T_1 initiated at site S_1 and creating agent at S_2,
- T_2 initiated at site S_2 and creating agent at S_3,
- T_3 initiated at site S_3 and creating agent at S_1.

Time	S_1	S_2	S_3
t_1	read_lock(T_1, x_1)	write_lock(T_2, y_2)	read_lock(T_3, z_3)
t_2	write_lock(T_1, y_1)	write_lock(T_2, z_2)	
t_3	write_lock(T_3, x_1)	write_lock(T_1, y_2)	write_lock(T_2, z_3)

Example - Distributed Deadlock

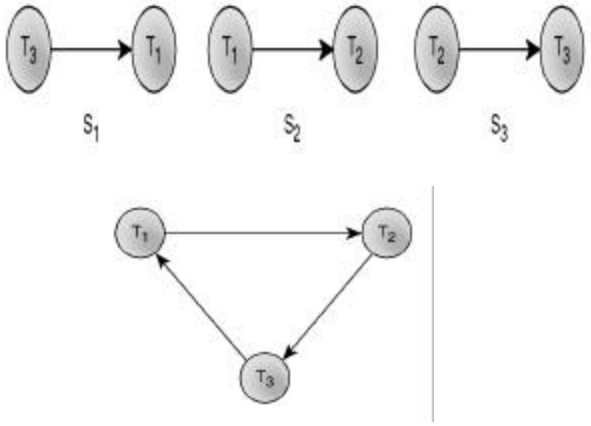

Centralized Deadlock Detection

- **Single site appointed deadlock detection coordinator (DDC).**
- **DDC has responsibility of constructing and maintaining GWFG.**
- **If one or more cycles exist, DDC must break each cycle by selecting transactions to be rolled back and restarted.**

Hierarchical Deadlock Detection

- **Sites are organized into a hierarchy.**
- **Each site sends its LWFG to detection site above it in hierarchy.**
- **Reduces dependence on centralized detection site.**

Hierarchical Deadlock Detection

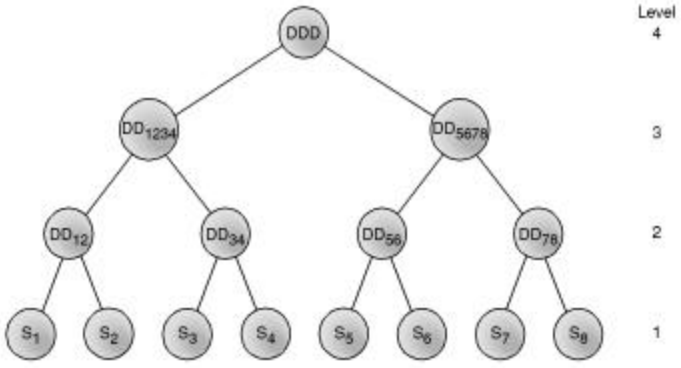

Two-Phase Commit (2PC)

- **Two phases: a *voting phase* and a *decision phase*.**
- **Coordinator asks all participants whether they are prepared to commit transaction.**
  - **If one participant votes abort, or fails to respond within a timeout period, coordinator instructs all participants to abort transaction.**
  - **If all vote commit, coordinator instructs all participants to commit.**
- **All participants must adopt global decision .**

Two-Phase Commit (2PC)

- **If participant votes abort, free to abort transaction immediately**
- **If participant votes commit, must wait for coordinator to broadcast global-commit or global-abort message.**
- **Protocol assumes each site has its own local log and can rollback or commit transaction reliably.**
- **If participant fails to vote, abort is assumed.**
- **If participant gets no vote instruction from coordinator, can abort.**

2PC Protocol for Participant Voting Commit

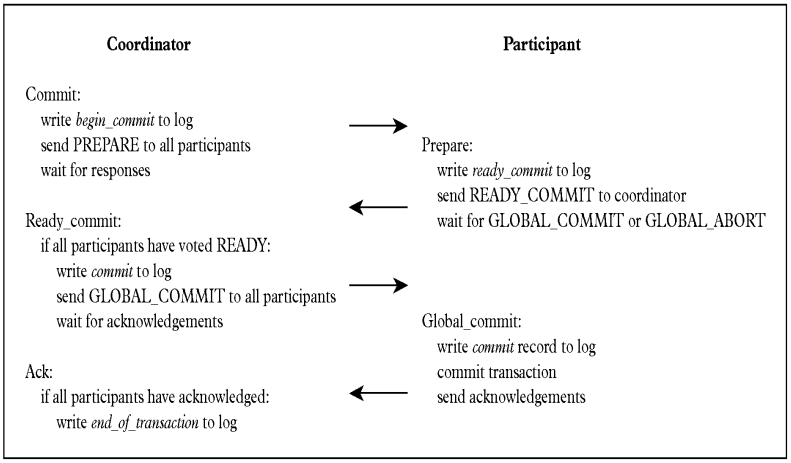

(a)

2PC Protocol for Participant Voting Abort

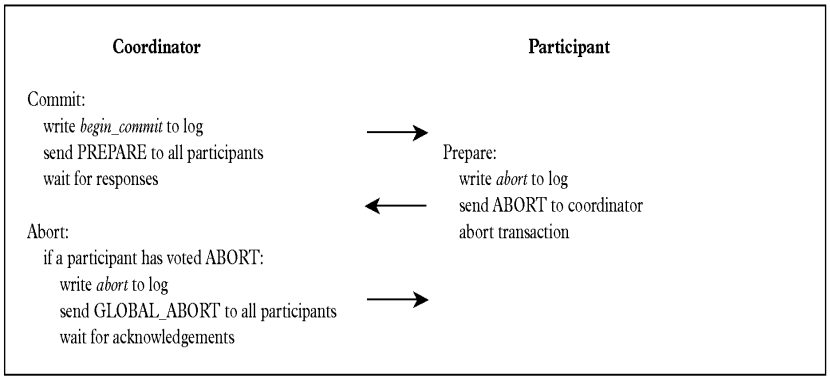

(b)

Termination Protocols

- **Invoked whenever a coordinator or participant fails to receive an expected message and times out.**

Coordinator

- **Timeout in WAITING state**
  - **Globally abort the transaction.**
- **Timeout in DECIDED state**
  - **Send global decision again to sites that have not acknowledged.**

Termination Protocols - Participant

- **Simplest termination protocol is to leave participant blocked until communication with the coordinator is re-established. Alternatively:**
- **Timeout in INITIAL state**
  - **Unilaterally abort the transaction.**
- **Timeout in the PREPARED state**
  - **Without more information, participant blocked.**
  - **Could get decision from another participant .**

State Transition Diagram for 2PC

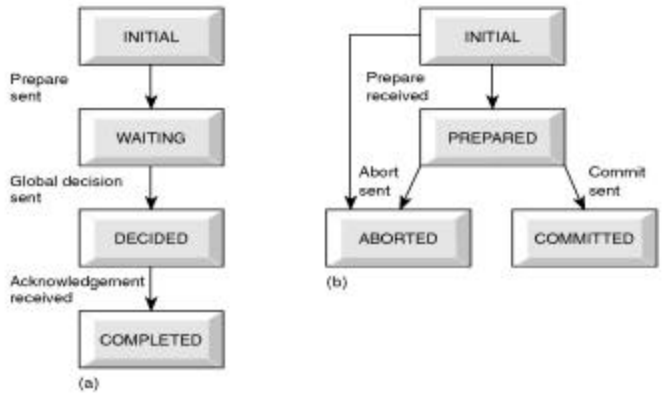

Recovery Protocols

- **Action to be taken by operational site in event of failure. Depends on what stage coordinator or participant had reached.**

Coordinator Failure

- **Failure in INITIAL state**
  - **Recovery starts the commit procedure.**
- **Failure in WAITING state**
  - **Recovery restarts the commit procedure.**

2PC - Coordinator Failure

- **Failure in DECIDED state**
  - **On restart, if coordinator has received all acknowledgements, it can complete successfully. Otherwise, has to initiate termination protocol discussed above.**

2PC - Participant Failure

- **Objective to ensure that participant on restart performs same action as all other participants and that this restart can be performed independently.**
- **Failure in INITIAL state**
  - **Unilaterally abort the transaction.**
- **Failure in PREPARED state**
  - **Recovery via termination protocol above.**
- **Failure in ABORTED/COMMITTED states**
  - **On restart, no further action is necessary.**

Three-Phase Commit (3PC)

- **2PC *is not* a non-blocking protocol.**
- **For example, a process that times out after voting commit, but before receiving global instruction, is blocked if it can communicate only with sites that do not know global decision.**
- **Probability of blocking occurring in practice is sufficiently rare that most existing systems use 2PC.**

Three-Phase Commit (3PC)

- **Alternative non-blocking protocol, called *three-phase commit (3PC)* protocol.**
- **Non-blocking for site failures, except in event of failure of all sites.**
- **Communication failures can result in different sites reaching different decisions, thereby violating atomicity of global transactions.**
- **3PC removes uncertainty period for participants who have voted commit and await global decision.**

Three-Phase Commit (3PC)

- **Introduces third phase, called *pre-commit,* between voting and global decision.**
- **On receiving all votes from participants, coordinator sends global pre-commit message.**
- **Participant who receives global pre-commit, knows all other participants have voted commit and that, in time, participant itself will definitely commit.**

State Transition Diagram for 3PC